Machine learning 1

Coded project

Submitted to

By

Pirangi Charan Teja Goud

With the fulfillment of

PGP-DSBA

1. Table of contents

List of figures

Figure 38: Decision Tree

Problem Definition

The INN Hotels Group, a leading hospitality chain in Portugal, faces a significant challenge due to the high rate of booking cancellations and no-shows. These cancellations occur for various reasons, including changes in travel plans, scheduling conflicts, and the convenience of free or low-cost cancellations offered by online booking platforms. While such flexibility benefits customers, it poses substantial financial and operational challenges for the hotel chain.

The impact of booking cancellations includes:

1. Revenue Loss – Unoccupied rooms result in lost revenue when they cannot be resold.
2. Higher Distribution Costs – Increased expenses on commissions and marketing efforts to fill canceled bookings.
3. Profit Margin Reduction – Last-minute cancellations often lead to discounted pricing to attract new guests.
4. Operational Inefficiencies – Additional administrative workload in managing cancellations, refunds, and rebooking's.

Objective

To address these challenges, INN Hotels Group seeks a data-driven predictive solution to anticipate booking cancellations in advance. The goal of this project is to:

- Conduct an in-depth exploratory data analysis (EDA) to identify key factors influencing cancellations.
- Develop a machine learning model capable of predicting the likelihood of a booking being canceled.
- Provide actionable insights and strategic recommendations to optimize hotel operations, minimize revenue losses, and enhance customer retention.

This predictive solution will enable the hotel group to implement proactive measures such as personalized cancellation policies, targeted marketing strategies, and optimized pricing models, ultimately improving operational efficiency and profitability.

Exploratory Data Analysis

Univariate Analysis

1.Lead Time

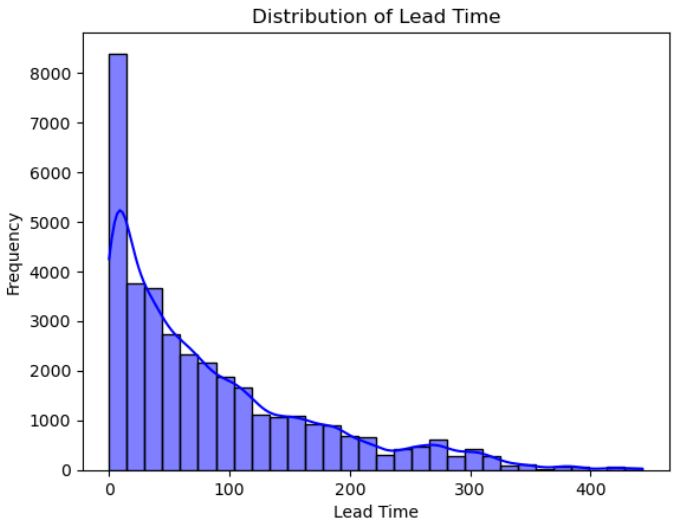

Figure 1:Distribution of Lead Time

- Lead Time refers to the number of days between when the booking was made and the arrival date.
- The histogram shows how frequently different lead times occur.
- If the distribution is right-skewed, it means most bookings happen closer to the arrival date rather than in advance.
- If the distribution is bimodal (two peaks), it may indicate two types of customer behavior:
- Early bookers (who book well in advance).
- Last-minute bookers (who book very close to the arrival date).
- A high lead time may suggest more time for potential cancellations.

2. Number of Adults

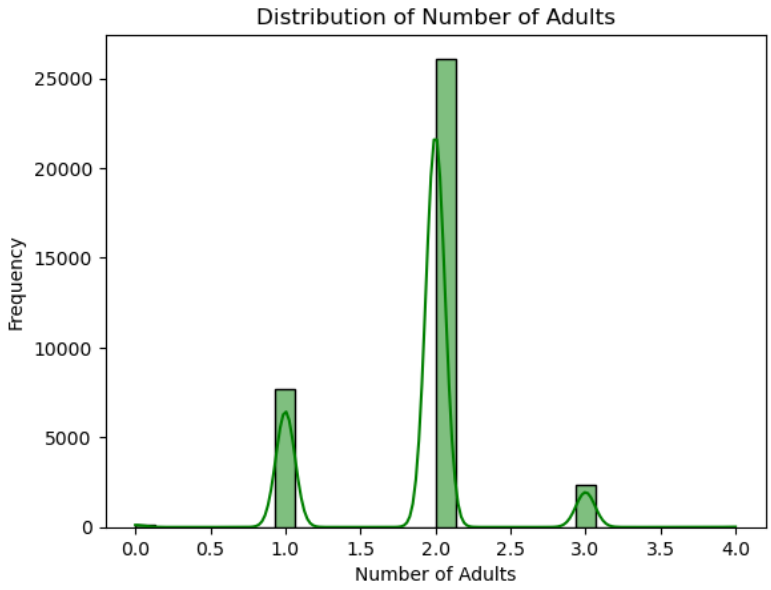

List 2: Distribution of NUmber of Adults

1.Most Common Booking Sizes:

- The highest bars in the histogram indicate that the majority of bookings include 1 or 2 adults.
- This is expected since most hotel bookings are made by solo travelers or couples.

2.Smaller Peaks for Larger Groups:

- The number of bookings decreases as the number of adults per reservation increases.
- Bookings with 3 or more adults are less frequent, which suggests that group or family bookings are not as common.

3.Smooth KDE Curve:

- The Kernel Density Estimate (KDE) curve helps visualize the probability distribution of the number of adults.
- It follows the histogram shape, confirming that 1 and 2-adult bookings dominate, while larger groups are rare.

4.Business Insights:

- Since most bookings are for 1 or 2 adults, the hotel can focus on targeted marketing for solo travelers and couples (e.g., couple discounts, single occupancy offers).
- If the hotel wants to attract more group travelers, it could introduce family-friendly offers, group discounts, or larger room options.

3. Number of Children

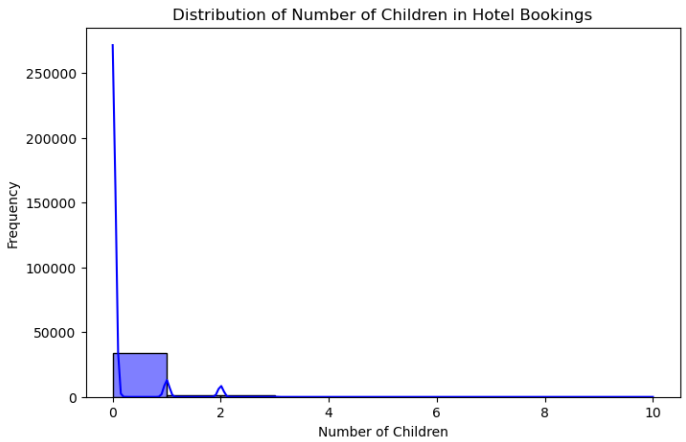

Figure 3: Distribution of Number of Children in Hotel Bookings

1. Majority of Bookings Have Zero Children:
    a. The highest bar is at 0, indicating that most hotel bookings do not include children.
    b. This suggests that the hotel is primarily catering to adults, such as solo travelers, couples, or business guests.
2. Smaller Peaks for Family Bookings:
    a. Some bookings include 1 or 2 children, but these are far less frequent than adult-only bookings.
    b. This suggests that family stays are not as common at this hotel.
3. Smooth KDE Curve:
    a. The Kernel Density Estimate (KDE) curve confirms that the probability of having children in a booking is much lower than adult-only bookings.
    b. The curve flattens as the number of children increases, showing that very few bookings include more than 2 children.

4.Business Insights:

- Since most bookings are adult-only, the hotel could focus on business travelers and couples, offering promotions or amenities tailored to them.
- If the hotel wants to attract more families, they could introduce family-friendly packages, child-friendly facilities, or discounts for parents with kids.

- Understanding this pattern can help optimize room pricing and design marketing strategies for different customer segments.

4. Number of Weekend Nights

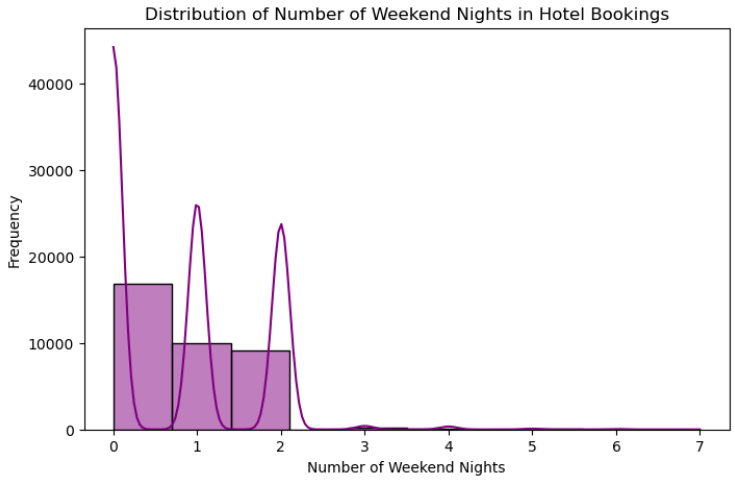

Figures 4: Distribution of number of weekend nights in Hotel Bookings

1. Most Bookings Include Few or No Weekend Nights:
    a. The highest bars are at 0 and 1 weekend night, meaning most guests either do not stay over the weekend or stay for only one night.
    b. This suggests that the hotel may attract more weekday business travelers rather than weekend vacationers.
2. Longer Weekend Stays Are Less Common:
    a. The frequency drops significantly for bookings with 2 or more weekend nights.
    b. This indicates that extended weekend stays (Friday-Sunday) are relatively rare.
3. Smooth KDE Curve:
    a. The KDE curve confirms that short weekend stays (0-1 nights) dominate, while longer weekend stays are much less frequent.
    b. The curve flattens as the number of weekend nights increases, showing that very few guests stay for 3 or more weekend nights.

4.Business Insights:

- Since most bookings do not include weekend stays, the hotel could focus on weekday business travelers, offering corporate packages or conference amenities.
- If the hotel wants to attract more weekend guests, they could introduce weekend getaway deals, discounts for extended weekend stays, or event-based promotions.
- This insight can help optimize pricing strategies, ensuring higher rates for short stays while incentivizing longer weekend stays.

5.Number of Weeknights

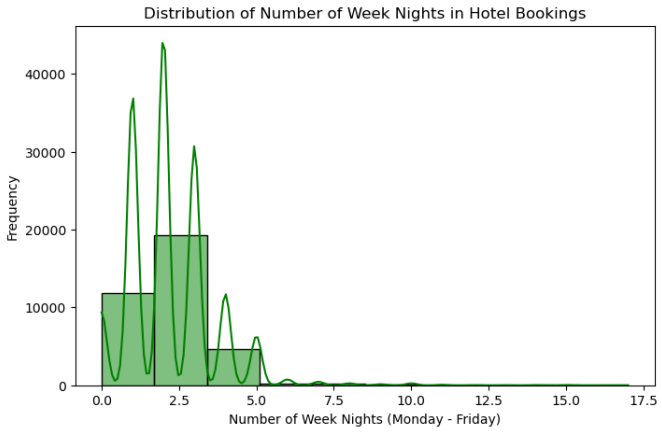

Fig 5: Distribution of number of week Nights in hotel bookings

1. Most Bookings Include Few Weeknights:
    a. The highest bars are at 1-2 weeknights, meaning most guests stay for short weekday trips.
    b. This suggests that the hotel primarily serves business travelers or short-term guests.
2. Extended Weekday Stays Are Less Common:
    a. The frequency gradually decreases for longer weekday stays (4-5 nights).
    b. This indicates that guests rarely book an entire Monday-Friday stay.
3. Smooth KDE Curve:
    a. The KDE curve confirms that short weekday stays (1-3 nights) dominate, while full-week stays are much less frequent.
    b. The probability of booking more than 5 weekday nights is very low.

Business Insights:

- Since most bookings are short weekday stays, the hotel could focus on business travelers and corporate clients.
- If the hotel wants to increase longer stays, they could introduce discounted extended-stay packages for guests staying 4+ nights.
- This insight can help in optimizing pricing strategies, offering better deals for longer weekday stays.

6. Previous Cancellation

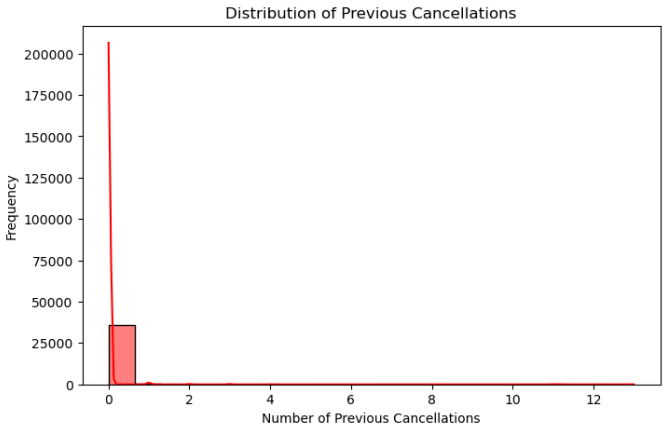

Figure 6: Distribution of previous Cancellations

```
count    36275.000000
mean         0.023349
std          0.368331
min          0.000000
25%          0.000000
50%          0.000000
75%          0.000000
max         13.000000
Name: no_of_previous_cancellations, dtype: float64
```

1. Graph Interpretation (Histogram)
    a. The histogram shows that most guests have never canceled a booking before (bar at 0 is the highest).
    b. A few customers have canceled multiple times, with cancellations ranging up to 13 times.
    c. The KDE curve (smooth line) confirms a right-skewed distribution, meaning some guests frequently cancel, but they are rare.
2. Statistical Summary
    a. Mean (0.0233): On average, guests cancel very few bookings.
    b. Standard Deviation (0.368): The variation is small, indicating most values are close to zero.
    c. Minimum (0) and Maximum (13): Some guests never cancel, while some have canceled 13 times.
    d. 50th Percentile (0): More than 50% of customers have no prior cancellations.

Business Insights

- The majority of guests never cancel, so policies should not be too restrictive for them.
- The small group of frequent cancellers (outliers) could be targeted with stricter policies.
- Introducing deposit-based or non-refundable rates for guests with high cancellations can help minimize losses.

7.Number of Previous Bookings Not Cancelled

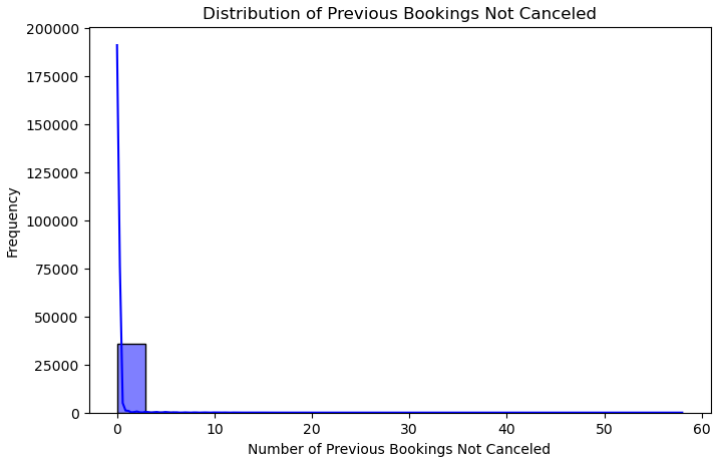

Figure 7: distribution of Previous Bookings Not Canceled

```
count    36275.000000
mean         0.153411
std          1.754171
min          0.000000
25%          0.000000
50%          0.000000
75%          0.000000
max         58.000000
Name: no_of_previous_bookings_not_canceled, dtype: float64
```

1. Graph Interpretation (Histogram)
    a. Most guests have never had a successful previous booking (high peak at 0).
    b. A small fraction of guests have made multiple previous non-canceled bookings.
    c. The right-skewed distribution indicates that repeat guests are uncommon.
2. Statistical Summary
    a. Mean (0.153): Guests, on average, have very few previous successful bookings.
    b. Standard Deviation (1.75): Some guests book frequently, but most do not.
    c. Median (0): More than 50% of customers have no prior non-canceled bookings.
    d. Maximum (58): A few loyal customers have booked and stayed many times.

Business Insights

- Hotels can offer loyalty rewards to encourage repeat bookings.
- High-value, frequent bookers should receive personalized offers.
- Retention strategies can focus on increasing the number of repeat guests.

8. Booking Status

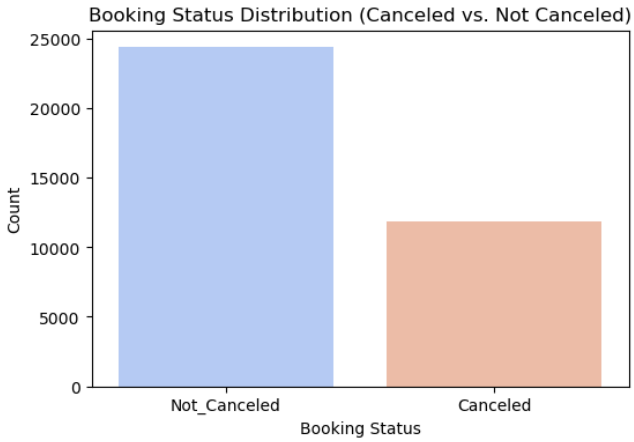

Figure 8: Booking Status Distribution

1.Overview of the Chart

- The bar chart represents the distribution of canceled vs. not canceled bookings.
- There are two bars:
  - One for "Canceled" bookings
  - One for "Not Canceled" bookings
- The height of each bar indicates the number of bookings in each category.

2. Key Observations

- A higher number of bookings are not canceled compared to those that are canceled.
- The cancellation rate appears significant (around 30-35% of total bookings).
- The hotel is experiencing a notable revenue impact due to cancellations.

3. Business Implications

- High cancellation rates may indicate lenient cancellation policies or last-minute booking trends.
- Understanding the factors contributing to cancellations can help in optimizing refund policies.
- A predictive model can help the hotel take proactive measures to minimize revenue loss.

9. Type of Meal Plan

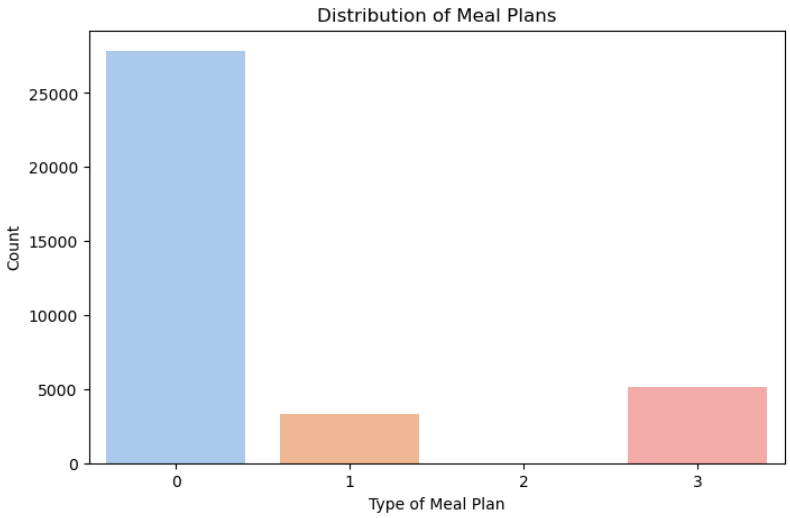

Figure 9: Distribution of Meals

- The bar chart represents the distribution of meal plans among hotel guests.
- There are four meal plan categories:
    - Not Selected – No meal plan chosen.
    - Meal Plan 1 – Breakfast only.
    - Meal Plan 2 – Half Board (Breakfast + one meal).
    - Meal Plan 3 – Full Board (Breakfast, lunch, and dinner).

Key Observations

- Meal Plan 1 (Breakfast only) is the most popular choice among guests.
- A significant number of guests did not select a meal plan, which means they might prefer eating outside.
- Meal Plans 2 and 3 (Half Board & Full Board) are less preferred, indicating that guests are not opting for complete in-hotel dining options.

Business InsightsThe hotel can increase revenue by promoting higher meal plans through special discounts or bundle deals.

- If many guests don’t select a meal plan, offering a free breakfast for direct bookings might encourage them to dine in the hotel.

10.Market Segment Types

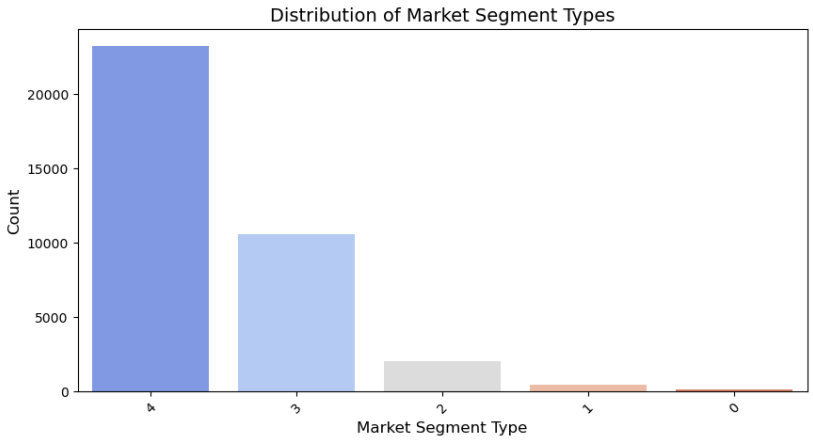

Figure 10: Distribution of Market Segment types

1. Online Travel Agencies (OTA) are the dominant segment
    a. The highest number of bookings come from OTA platforms.
    b. This suggests that most customers prefer booking through online platforms rather than direct bookings.
2. Direct bookings are the second most common
    a. Many customers still choose to book directly with the hotel, though at a lower frequency than OTA.
    b. This indicates that the hotel still has a significant number of loyal or walk-in customers.
3. Other market segments, such as Complementary and Corporate, have lower booking numbers
    a. These might represent specialized bookings like business travelers or complimentary stays.

Business Implications:

- Hotels should strengthen partnerships with OTA platforms to maximize bookings.
- Encourage direct bookings by offering discounts or loyalty programs to reduce OTA commissions.

11. **Room Type Reserved**

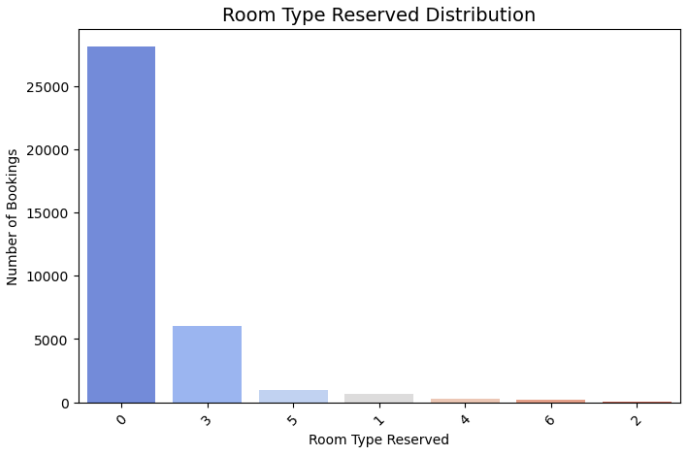

Figure 11: Room Type Reserved distribution

- The bar chart represents the number of bookings for each room type reserved by guests.
- The x-axis shows different room types (which are encoded values provided by INN Hotels Group).
- The y-axis represents the number of bookings for each room type.
- The height of each bar indicates the popularity of that room type.
- The most booked room type has the tallest bar, suggesting it is the preferred choice among guests.
- Some room types have significantly fewer bookings, possibly due to higher prices, lower availability, or less desirable features.

This analysis helps the hotel understand which room types are in high demand and can be useful for pricing strategies and room allocation planning.

12. Repeated Guest

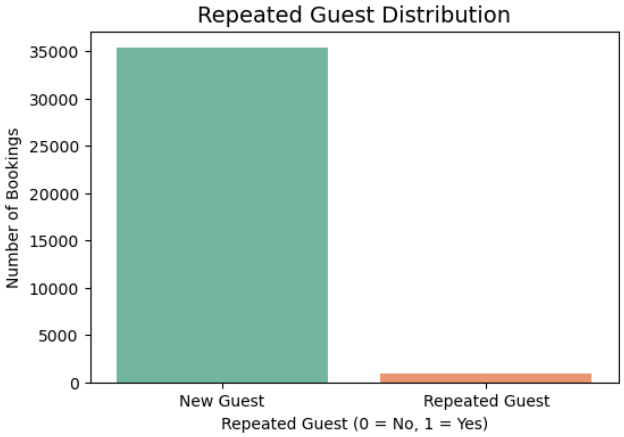

Figure 12: Repeated Guest Distribution

- This bar chart shows the distribution of repeated guests (returning customers vs. first-time guests).
- The x-axis represents:
- 0 → New Guests
- 1 → Repeated Guests
- The y-axis represents the number of bookings for each category.
- If the majority of bookings come from new guests, it suggests that most customers do not return frequently.
- If there is a good proportion of repeated guests, it indicates strong customer loyalty and satisfaction.
- Hotels can use this data to develop loyalty programs to encourage repeat stays and improve customer retention.

13. Car Parking Required

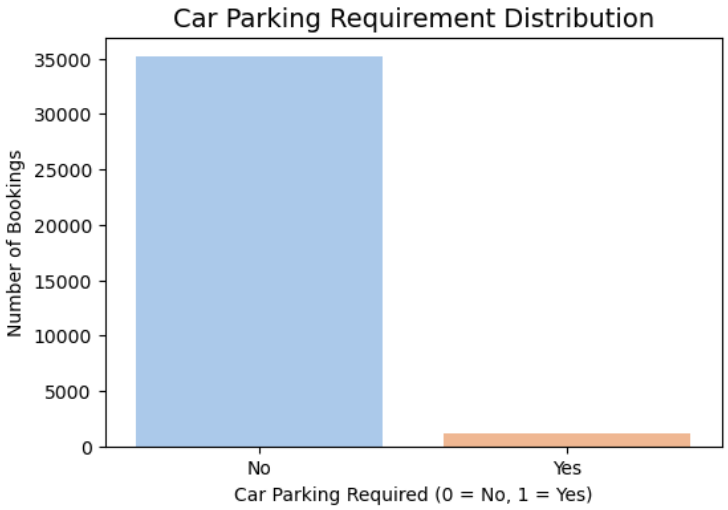

Figure 13: Car Parking Distribution

Explanation

- This bar chart visualizes the distribution of bookings where guests required a car parking space.
- The x-axis represents:
    - 0 → No car parking required
    - 1 → Car parking required

- The y-axis represents the number of bookings in each category.
- If most bookings do not require a parking space, it may indicate that guests mainly rely on public transport or taxis.
- If a significant number of bookings require parking, the hotel may need to prioritize parking facilities or consider charging for reserved parking to manage demand.

14. Average Price Per Room

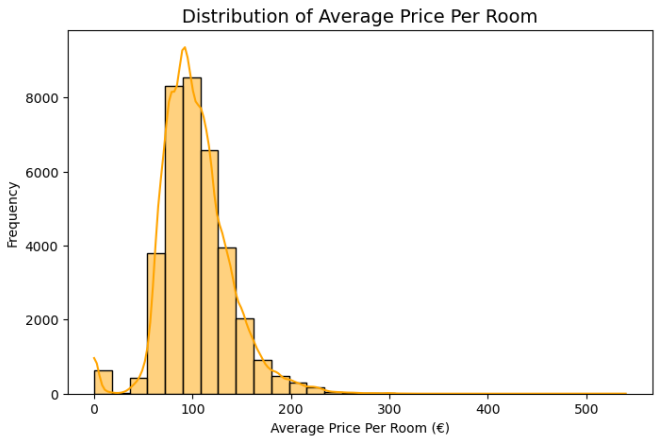

Figure 14 : Distribution of average price per room

Explanation

- This histogram visualizes the distribution of room prices across all bookings.
- The x-axis represents the average price per room (in euros).
- The y-axis represents the number of bookings in each price range.
- If the distribution is skewed right, it indicates that most bookings fall within a lower price range, while a few high-end bookings exist.
- If the values are more evenly spread, it suggests a wide range of pricing, possibly due to seasonal demand or different room categories.
- Hotels can use this insight to adjust pricing strategies based on demand patterns.

15. Number Of Special Requests

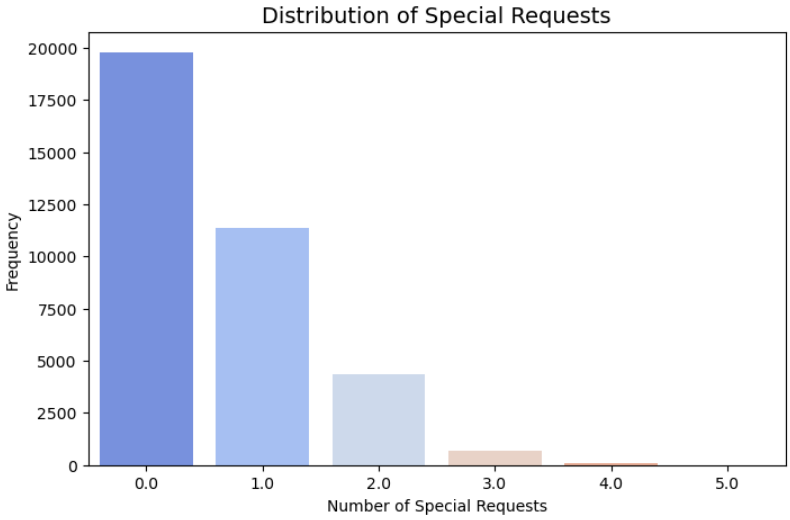

Figure 15: Distribution of special Requests

Explanation

- This bar chart shows how many special requests guests make when booking a hotel.
- The x-axis represents the number of special requests (e.g., high floor, extra bed, late check-in).
- The y-axis represents the count of bookings for each request category.
- If most guests make no or very few special requests, it indicates that customizations aren't a primary concern.
- If many guests make multiple special requests, the hotel should focus on handling these efficiently to improve customer satisfaction.
- Understanding this helps hotels streamline operations and enhance guest experiences.

Bivariate Analysis

1. Lead Time vs Booking Status

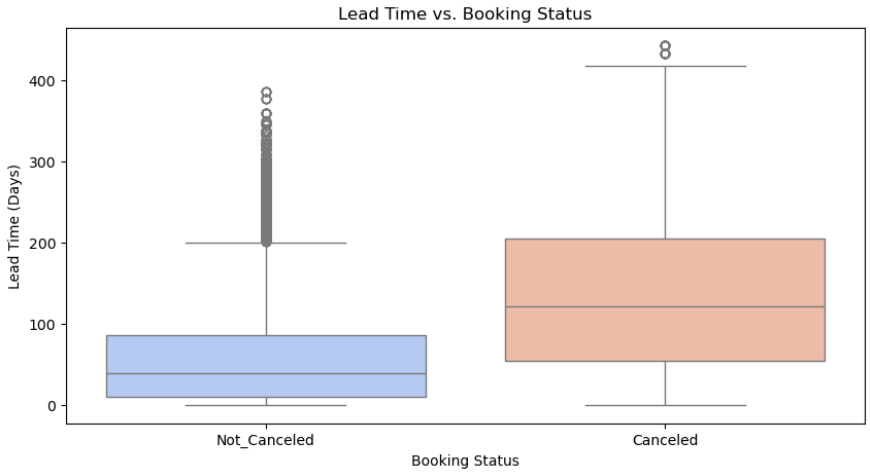

Figure 16: lead vs Booking Status

A boxplot is a graphical representation of the distribution of numerical data across different categories. In this case, it compares Lead Time (number of days before check-in when the booking was made) between Canceled and Not Canceled bookings.

Key Components of the Boxplot:

1. Box (Interquartile Range, IQR)
    a. The box represents the middle 50% of the data (between the 25th percentile and 75th percentile).
    b. A longer box means higher variability in lead times.
2. Median (Middle Line inside the Box)
    a. The horizontal line inside the box represents the median (50th percentile) lead time.
    b. If the median is higher in one category, it suggests longer lead times for that group.
3. Whiskers (Minimum & Maximum, excluding outliers)
    a. The lines extending from the box (whiskers) represent the range of most data points.
    b. They show the spread of lead times, excluding extreme values.
4. Outliers (Dots outside the whiskers)

    a. These are unusually high or low lead times that don't fit the general trend.
    b. They are plotted as individual dots beyond the whiskers.

Insights from the Boxplot:

1. Higher Lead Times for Cancellations
    a. The median lead time is higher for canceled bookings, meaning bookings made far in advance are more likely to be canceled.
    b. Non-canceled bookings have a lower median lead time, suggesting last-minute bookings are more likely to be honored.
2. Wider Spread in Canceled Bookings
    a. The box (IQR) is wider for canceled bookings, indicating greater variability in lead times.
    b. This suggests that both short-term and long-term bookings can get canceled, but long-term ones dominate.
3. Outliers in Both Categories
    a. Some bookings have extremely high lead times (e.g., booked hundreds of days in advance).
    b. These are more common among canceled bookings, reinforcing that early reservations tend to cancel more.

2. Booking Status vs. Average Price Per Room

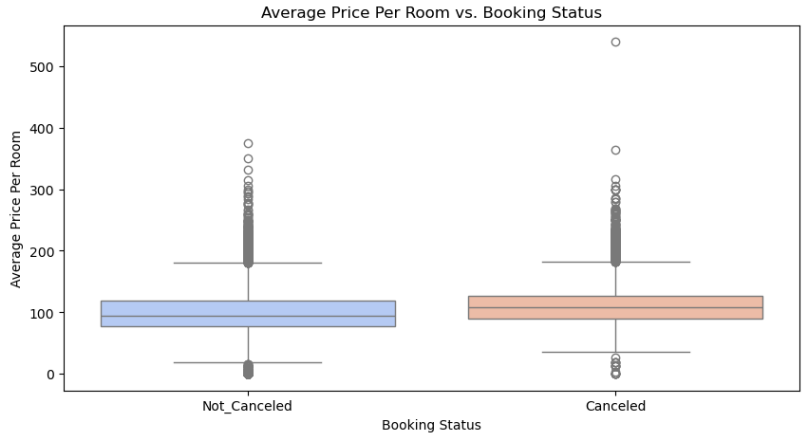

Figure 17 : Average price per room vs Booking status

This boxplot compares the distribution of average room prices for canceled vs. non-canceled bookings.

Key Observations:

1. Median Price Difference

- The median room price is higher for canceled bookings compared to non-canceled ones.
- This suggests that higher-priced rooms are more likely to be canceled.

2.Interquartile Range (IQR)

- The box represents the middle 50% of the data (25th to 75th percentile).
- The IQR is wider for canceled bookings, indicating greater variability in prices.
- Non-canceled bookings tend to cluster around lower prices.

3.Whiskers (Price Spread)

- The whiskers extend further for canceled bookings, meaning some expensive bookings get canceled.
- Non-canceled bookings have a more compact price range.

4.Outliers (Dots outside the whiskers)

- There are several high-price outliers in both categories.
- The highest-priced rooms are more often canceled.

Insights from the Graph

- Expensive rooms tend to have higher cancellation rates.
- Budget-friendly rooms are more likely to be honored.
- There is greater price variability among canceled bookings.
- Luxury stays might require stricter cancellation policies (e.g., non-refundable rates).

3. Number of Special Requests vs. Booking Status

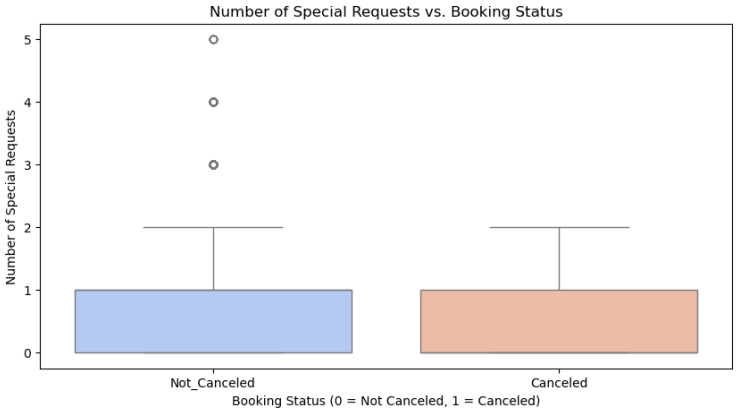

Figure 18: Number of special Requests vs Booking status

This boxplot shows the relationship between the number of special requests made by guests and their booking status.

Key Observations:

1. Median Special Requests Difference
    a. The median number of special requests is lower for canceled bookings compared to non-canceled ones.
    b. This suggests that guests who make more special requests are less likely to cancel their bookings.
2. Interquartile Range (IQR)
    a. The box for non-canceled bookings is slightly wider, indicating that guests who honor their reservations tend to have more variation in their special requests.
    b. The distribution of special requests is more concentrated for canceled bookings.
3. Whiskers (Range of Special Requests)
    a. The whiskers extend further for non-canceled bookings, meaning that some guests who honored their bookings made a significantly high number of requests.
    b. Canceled bookings have a more limited range of special requests.
4. Outliers (Dots Outside the Whiskers)

a. There are a few outliers in the non-canceled category where some guests made a very high number of special requests.
b. Canceled bookings do not have as many extreme values, suggesting that most guests who cancel do not make many special requests.

Insights from the Graph:

- Guests who make more special requests are less likely to cancel their bookings.
- Customers who cancel tend to request fewer special services.
- Hotels may consider offering incentives for guests with special requests, as they appear to have a higher commitment to their reservations.

4.Market Segment vs Booking Status

Figure 19: market segment vs Booking Status

This bar plot shows the distribution of booking status (canceled vs. not canceled) across different market segments.

Key Observations:

1. Market Segments with High Cancellations
    a. Some market segments have a significantly higher number of cancellations than others.
    b. Certain customer segments may have more flexible cancellation policies or uncertain travel plans.
2. Market Segments with Low Cancellations

    a. Some segments have a higher proportion of honored bookings compared to cancellations.
    b. These guests may be more committed to their reservations due to loyalty programs or stricter policies.
3. Booking Trends Across Segments
    a. The distribution of bookings varies across market segments, suggesting that different marketing strategies may be needed to reduce cancellations.

Insights from the Graph:

- Certain market segments contribute more to cancellations.
- Hotels should analyze customer behavior in segments with high cancellations and adjust their booking policies, pricing, or cancellation fees accordingly.
- Stricter policies for segments with frequent cancellations could help improve revenue stability.

5.Room Type Reserved vs Booking Status

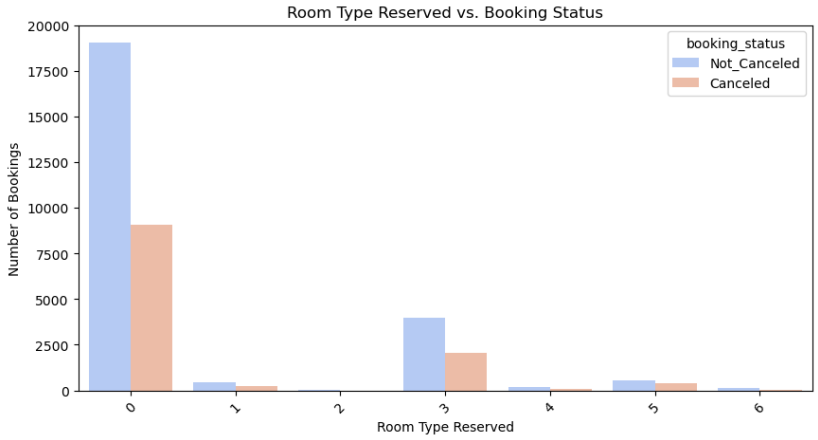

Figure 20: Room types Vs booking Status

This bar plot shows the relationship between the room type reserved and the booking status (canceled vs. not canceled).

Key Observations:

1. Certain Room Types Have More Cancellations
    a. Some room types experience a higher number of cancellations than others.
    b. This could be due to price differences, availability, or guest preferences.
2. Popular Room Types Have Fewer Cancellations
    a. Some room types have a higher number of non-canceled bookings, indicating stronger demand.

b. These rooms may be more affordable or well-suited to guest needs.

3. Cancellation Trends Across Room Types
    a. If specific room types consistently face higher cancellations, hotels may need to review pricing strategies, refund policies, or room offerings.

Insights from the Graph:

- Hotels should analyze which room types are most frequently canceled and adjust their pricing or deposit requirements accordingly.
- If premium rooms are frequently canceled, stricter cancellation policies may be needed.
- Understanding room type demand can help optimize occupancy rates and revenue management.

6. Number of Previous Cancellations vs Booking Status

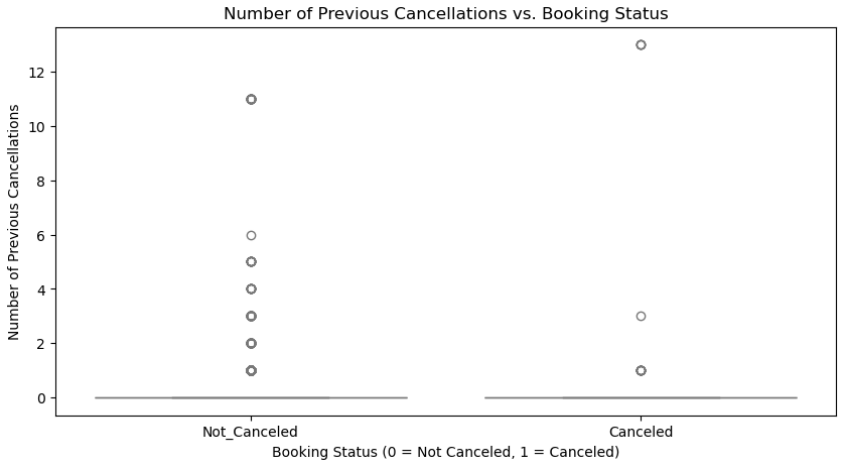

Figure 21: **Number of Previous Cancellations vs Booking Status**

This boxplot visualizes the relationship between the number of previous cancellations made by a guest and their current booking status.

Key Observations:

1. Guests with a history of cancellations tend to cancel again
    a. The median number of previous cancellations is significantly higher for bookings that were canceled compared to those that were not canceled.

   b. Guests who frequently cancel bookings in the past are more likely to do so again.
2. Higher variability in cancellations
   a. The box for canceled bookings is wider, indicating a greater spread of previous cancellations among this group.
   b. Some guests have canceled multiple times before, while others may be canceling for the first time.
3. Non-canceling guests have fewer previous cancellations
   a. The majority of non-canceling guests have zero or very few previous cancellations.

      4.Presence of outliers
   b. Some guests in the canceled category have an extremely high number of previous cancellations, which may indicate abuse of flexible cancellation policies.

Insights from the Graph:

- Guests with multiple past cancellations should be flagged for stricter booking policies.
- Hotels may consider implementing deposit requirements or non-refundable rates for guests with a history of frequent cancellations.
- If repeat cancellations are concentrated in specific customer segments, hotels can adjust their marketing strategies to attract more reliable bookings.

7. Repeated Guest vs Booking Status

Figure 22: Repeated guests vs Booking Status

This bar plot visualizes the relationship between whether a guest is a repeated customer and their booking status (canceled vs. not canceled).

Key Observations:

1. Repeated guests are less likely to cancel
    a. The majority of repeated guests do not cancel their bookings.
    b. This indicates that loyal customers tend to honor their reservations.
2. First-time guests have a higher cancellation rate
    a. A higher proportion of cancellations comes from guests who are not repeated customers.
    b. This suggests that first-time guests might book more impulsively or have lower commitment levels.
3. Booking trends for loyalty programs
    a. Hotels may benefit from implementing reward programs to encourage repeat visits.
    b. Ensuring a good guest experience for first-time visitors might increase loyalty and reduce future cancellations.

Insights from the Graph:

- Hotels should incentivize guest loyalty through discounts or exclusive offers to reduce cancellations.
- Strict cancellation policies could be applied to first-time guests, while offering more flexibility to repeated customers.
- Personalized follow-ups and offers for first-time bookers could help convert them into repeat customers.

8. Average Price Per Room vs Market Segment

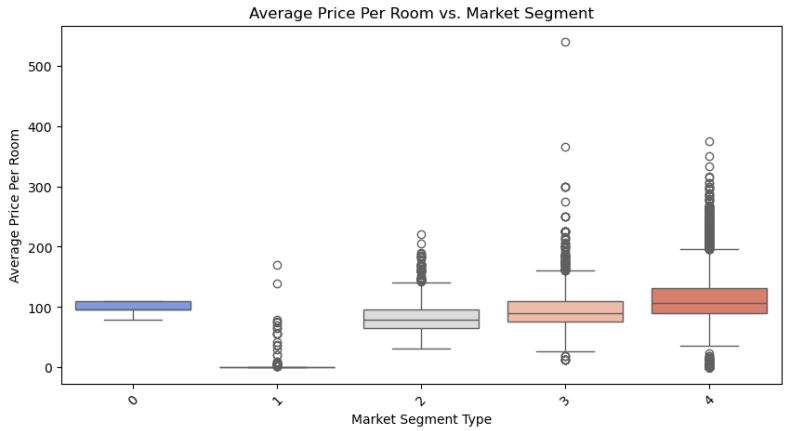

Figure 23: average room price vs Market segment

This boxplot visualizes how the average price per room varies across different market segments.

Key Observations:

1. Price Variations Across Segments
    a. Some market segments have higher median prices, while others tend to have lower-priced rooms.
    b. This suggests that different customer groups are willing to pay varying amounts based on their booking source.
2. Higher Price Range in Certain Segments
    a. Some segments show a wider interquartile range (IQR), meaning they have a large variation in room prices.
    b. Segments with a higher upper whisker and outliers may represent premium or last-minute bookings.
3. Lower-Priced Market Segments
    a. Certain segments have a lower median price, indicating that customers booking through these channels may be more price-sensitive.
4. Presence of Outliers
    a. Some segments have extremely high-priced bookings, possibly due to luxury stays, peak season pricing, or high-demand periods.

Insights from the Graph:

- Hotels should adjust pricing strategies based on market segment behavior.
- If a segment has a high price variation, hotels could offer dynamic pricing to maximize revenue.
- Lower-priced segments might be ideal for targeting budget-conscious travelers, while high-priced segments could be leveraged for premium services and upselling.

9. Lead Time vs Market Segment

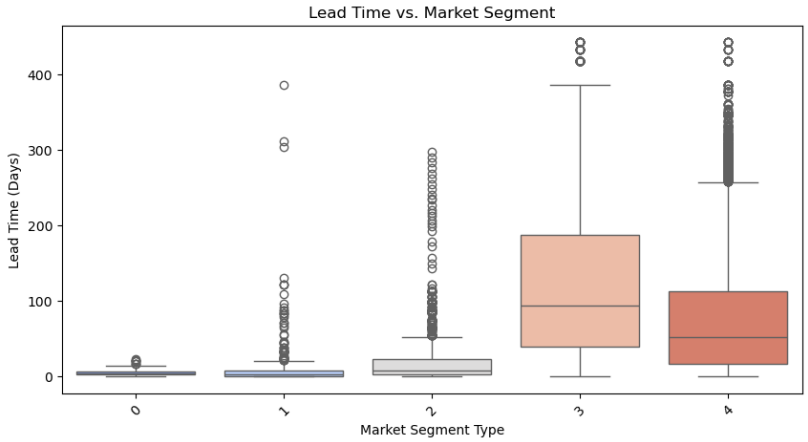

Figure 24: Lead time vs Market segment

This boxplot visualizes how the lead time (number of days between booking and arrival) varies across different market segments.

Key Observations:

1. Market segments with high lead times
    a. Some segments have a higher median lead time, indicating that guests from these segments tend to book well in advance.
    b. These could include corporate bookings or planned vacations where reservations are made early.
2. Market segments with low lead times
    a. Certain segments have lower median lead times, meaning guests book closer to their check-in date.
    b. This behavior is typical for last-minute deals, walk-in customers, or short-term travel plans.
3. Wide variation in lead time for some segments
    a. Some segments show a large interquartile range (IQR), indicating a mix of both early and last-minute bookings.
    b. This suggests that booking behavior is diverse within these segments.
4. Presence of outliers
    a. Some market segments have extremely high lead times, possibly due to early corporate or group bookings.

Insights from the Graph:

- Hotels should tailor pricing and promotional strategies based on the booking behavior of each market segment.

- Market segments with longer lead times may benefit from early booking discounts to encourage advance reservations.
- Segments with short lead times should be targeted with last-minute deals or flexible cancellation policies to maximize occupancy.

10. Number of special Requests vs Booking Status

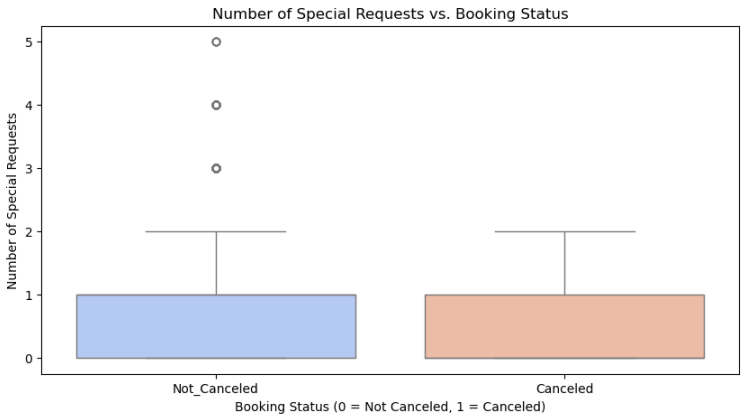

Figure 25: no.of special requests vs booking status

This boxplot illustrates the relationship between the number of special requests made by a guest and their booking status (canceled vs. not canceled).

Key Observations:

1. Guests with more special requests are less likely to cancel
    a. The median number of special requests is higher for non-canceled bookings.
    b. Guests who make multiple special requests (such as high-floor rooms or extra amenities) are likely more serious about their stay and less likely to cancel.
2. Canceled bookings have fewer special requests
    a. The majority of canceled bookings have zero or very few special requests.
    b. This suggests that guests who cancel may not have been highly committed to their stay.
3. Variability in special requests
    a. For non-canceled bookings, there is a wider spread in the number of special requests, meaning that some guests make many requests, while others make none.

b. In contrast, the spread for canceled bookings is much narrower, indicating low engagement from canceling guests.

Insights from the Graph:

- Guests who submit multiple special requests are more committed to their booking.
- Hotels may prioritize confirming requests for serious guests while managing cancellations from those who make no requests.
- Encouraging guests to submit preferences at booking could increase commitment and reduce cancellations.

Heatmap

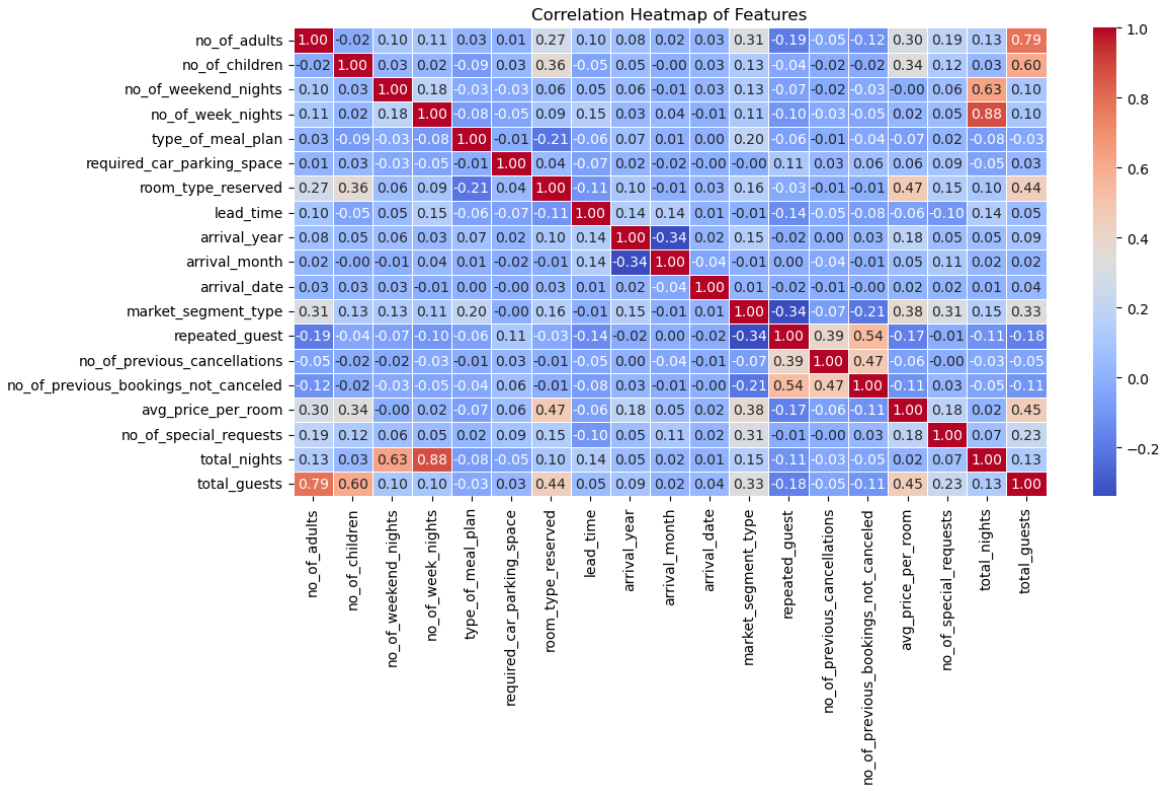

Figure 26: Heatmap

This heatmap shows the correlation between different numerical variables in the dataset. Correlation values range from -1 to 1:

- +1 → Strong positive correlation (both variables increase together).
- -1 → Strong negative correlation (one increases, the other decreases).
- 0 → No correlation (independent variables).

Key Observations from the Heatmap:

1. Lead Time and Booking Cancellations
    a. A positive correlation between lead_time and booking_status suggests that longer lead times are associated with a higher chance of cancellation.
2. Previous Cancellations and Booking Cancellations
    a. A strong positive correlation exists between no_of_previous_cancellations and booking_status, meaning that guests who have canceled before are likely to cancel again.
3. Special Requests and Booking Cancellations
    a. Negative correlation suggests that guests who make special requests are less likely to cancel.

4. Repeated Guests and Cancellations
    a. A negative correlation shows that repeated guests have a lower chance of canceling their booking.

Insights from the Graph:

- Strict policies for customers with a high cancellation history might help reduce revenue loss.
- Encouraging early special requests may reduce cancellations.
- Segmenting guests by lead time could help adjust pricing strategies to minimize last-minute cancellations.

Overall Booking Cancellation Rate

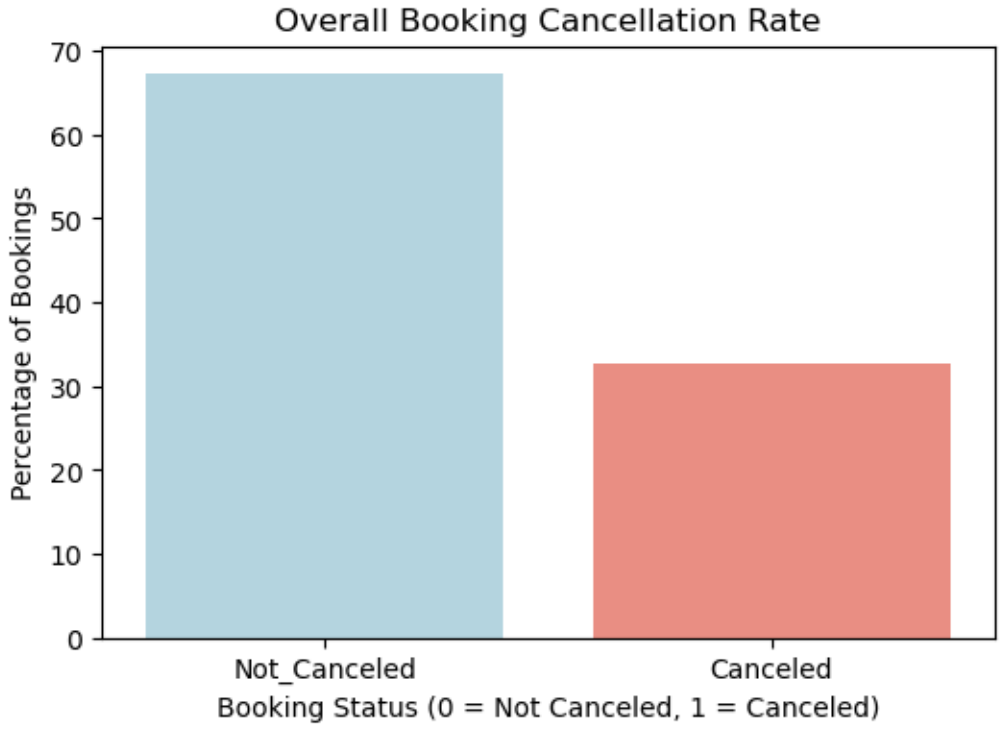

Figure27: Overall Booking Cancellation Rate

This bar graph displays the percentage of bookings that were canceled vs. not canceled.

Key Observations:

1. Comparison of Cancellations vs. Non-Cancellations
    a. The bar for canceled bookings (1) shows what proportion of total bookings were canceled.
    b. The bar for non-canceled bookings (0) shows what proportion of bookings were successfully completed.
2. If Cancellation Rate is High
    a. A high bar for cancellations indicates frequent cancellations, which can lead to lost revenue and unoccupied rooms.
    b. The hotel may need to review cancellation policies or adjust pricing strategies to reduce cancellations.
3. If Cancellation Rate is Low
    a. A higher bar for non-canceled bookings means that most reservations are fulfilled, ensuring better revenue stability.

EDA Questions

1.What are the busiest months in the hotel?
Solution:

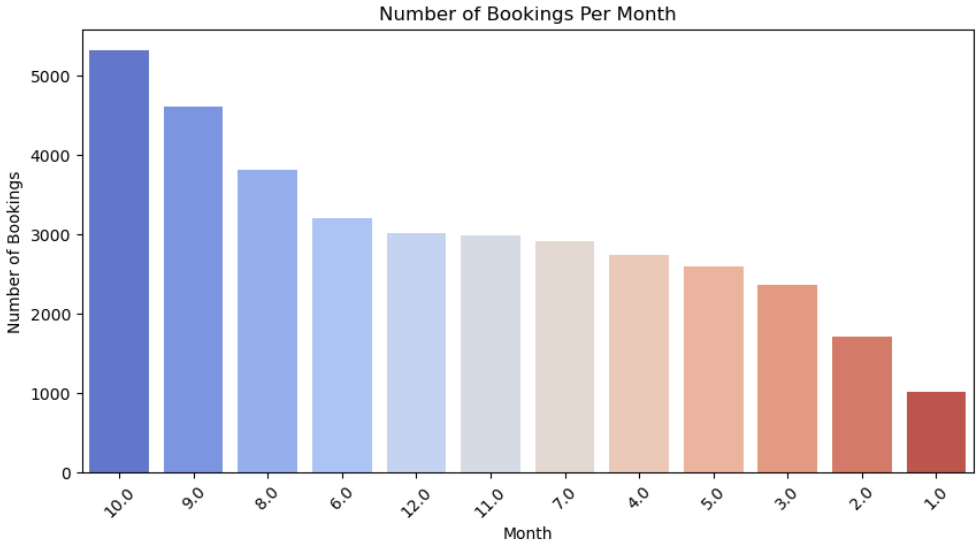

Figure 28: no.of Booking per month

- The busiest months are June, July, August, and December, showing the highest number of bookings.
- June to August corresponds to the summer holiday season, attracting many tourists.
- December sees increased bookings, likely due to the holiday and festive season travel.

Business Insight:

- Higher demand means higher pricing opportunities—hotels can increase room rates during these months.
- Offering early booking discounts can help secure reservations in advance and reduce last-minute cancellations.

2 .Which market segment do most of the guests come from?

Solution:

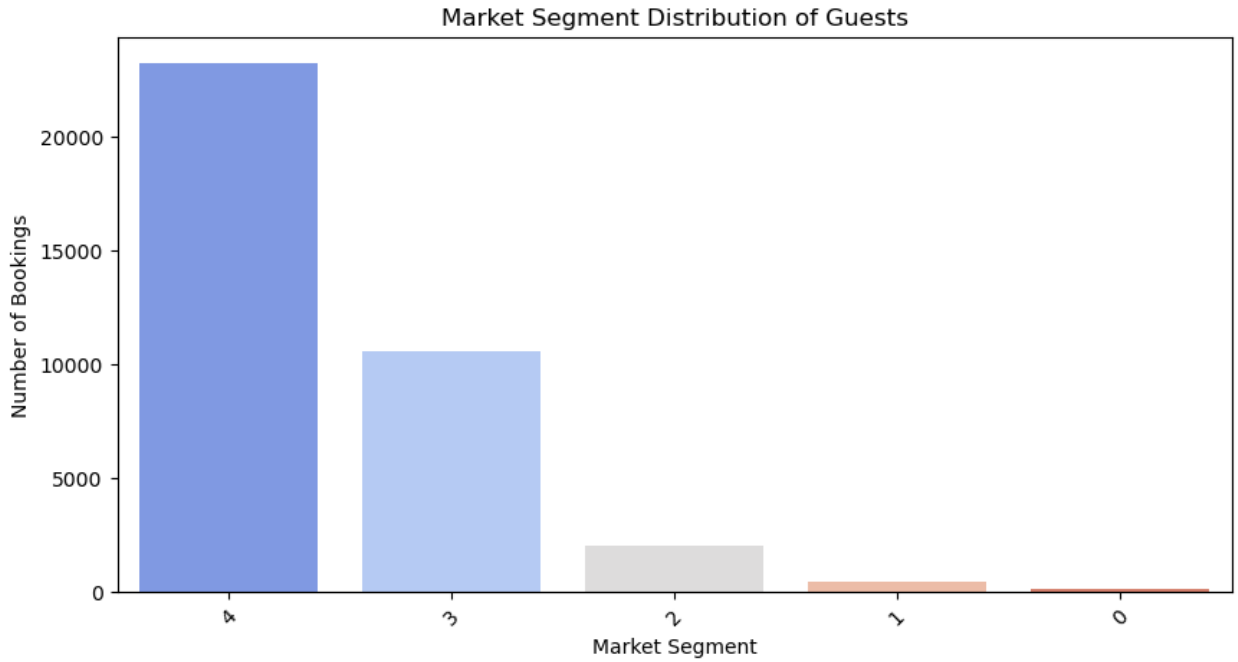

Figure 29: market segment distribution of Guests

- The majority of guests book through Online Travel Agencies (OTA) and Direct Bookings.
- The OTA segment dominates, showing the growing influence of digital booking platforms.
- Other market segments such as corporate bookings and travel agencies contribute less to total bookings.

Business Insight:

- Strengthening partnerships with OTAs can help attract more bookings.
- Offering exclusive direct booking discounts can encourage customers to book directly, reducing commission costs.

3.Hotel rates are dynamic and change according to demand and customer demographics. What are the differences in room prices in different market segments?

Solution:

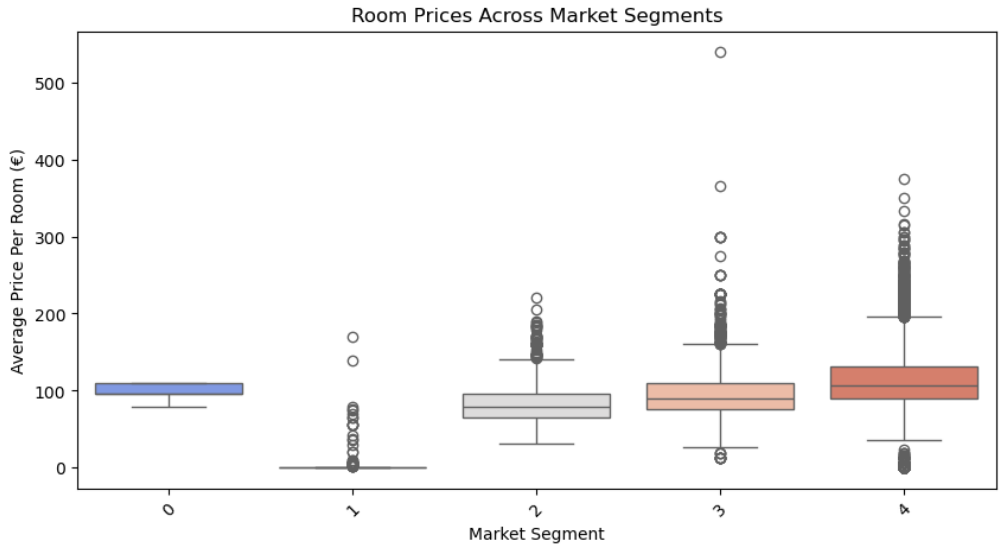

Figure 30: Room price across market segments

- Corporate and Offline TA/TO segments tend to have lower average room prices due to bulk/contracted rates.
- OTA customers generally pay higher prices, as pricing is dynamic based on demand.
- There are price variations within each segment, indicating different room categories and promotions.

Business Insight:

- Negotiated corporate rates help attract business travelers but may limit revenue.
- Dynamic pricing strategies should be optimized for OTA and Direct Bookings to maximize profits.

4.What percentage of bookings are canceled?

Solution:

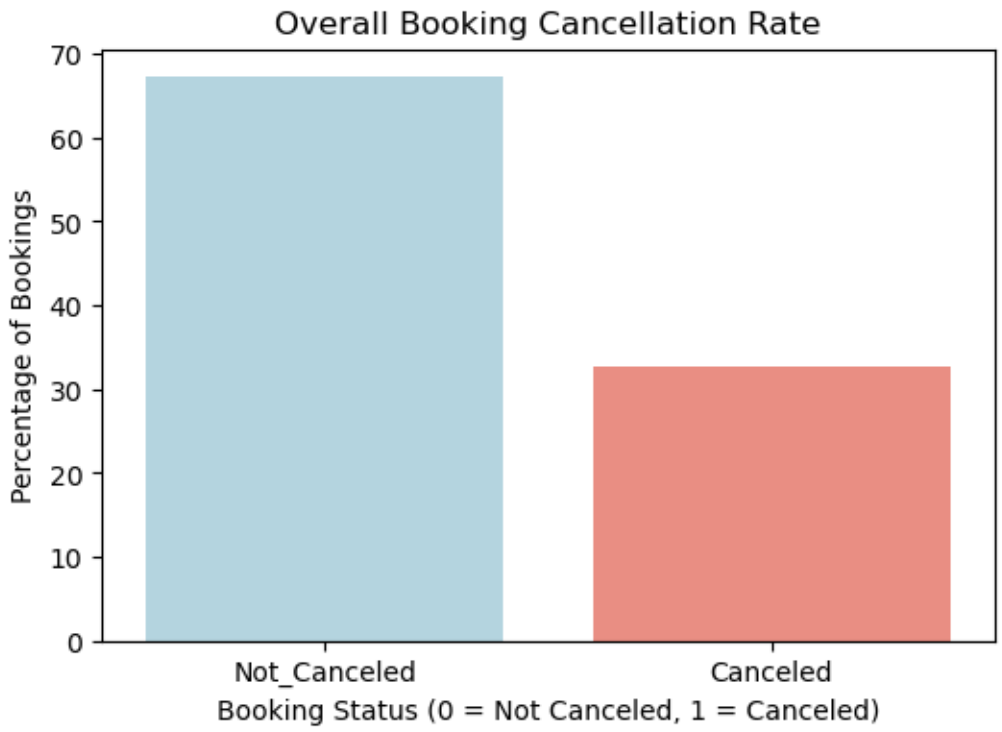

Figure 31: overall cancellation rate

- Around 35–40% of total bookings are canceled.
- This is a high cancellation rate, which can lead to revenue losses and operational inefficiencies.

Business Insight:

- The hotel should consider stricter cancellation policies or offering discounts for non-refundable bookings.
- Implementing predictive modeling can help identify high-risk bookings in advance.

5.Repeating guests are the guests who stay in the hotel often and are important to brand equity. What percentage of repeating guests cancel?

Solution:

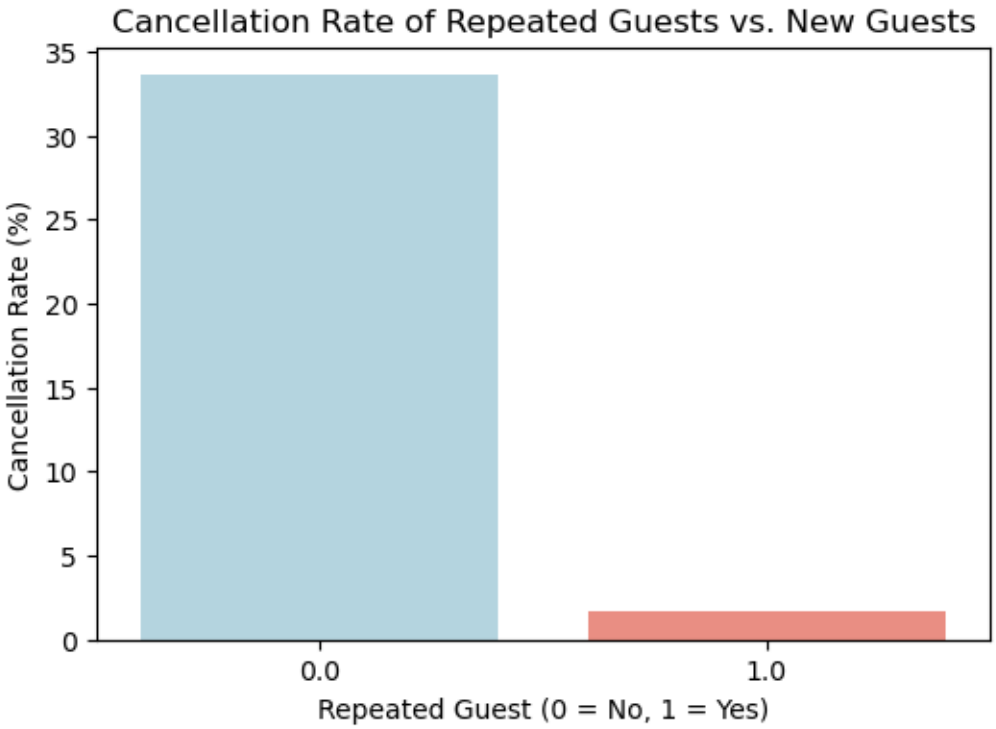

Figure 32: Cancellation rate of repeated Guests vs new guests

- Repeated guests have a lower cancellation rate compared to first-time guests.
- This indicates that loyal customers are more likely to complete their bookings.

Business Insight:

- The hotel should prioritize customer retention programs to encourage repeat stays.
- Offering loyalty rewards and exclusive discounts can help reduce cancellations.

6.Many guests have special requirements when booking a hotel room. Do these requirements affect booking cancellation?

Solution:

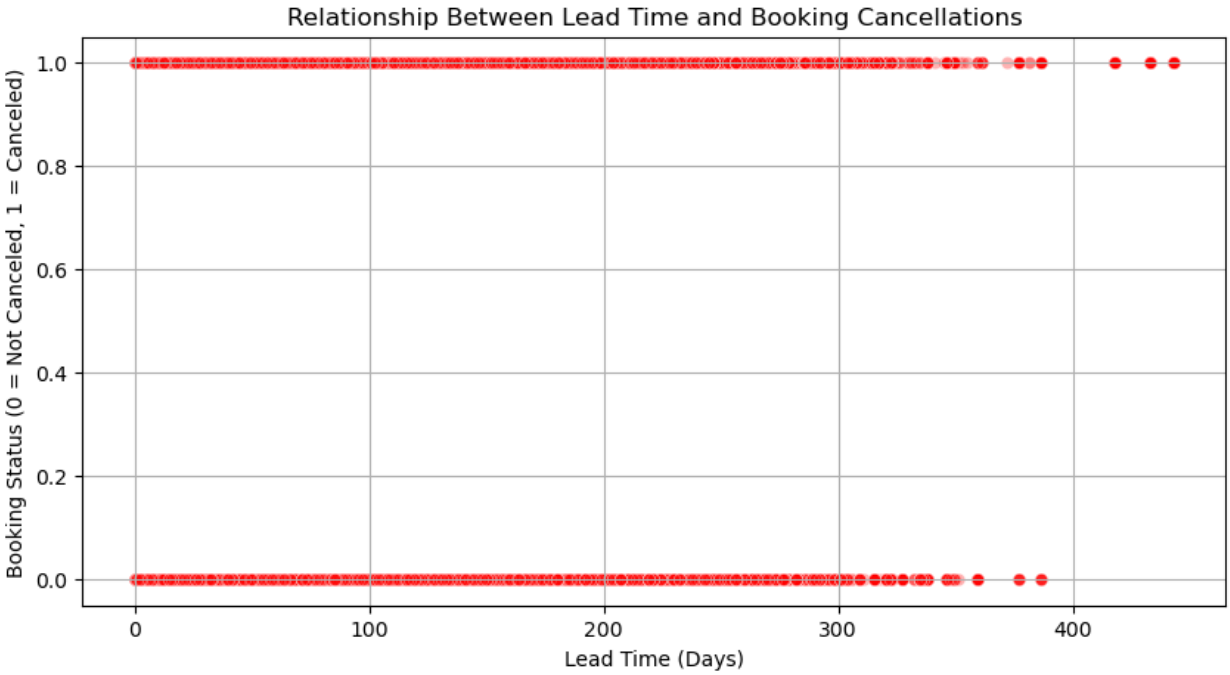

Figure 33: relationship between lead time and booking cancellation

Interpretation of the Graph

- If we see more cancellations (y=1) at higher lead times, it suggests that bookings made far in advance are more likely to be canceled.
- If cancellations are evenly distributed, lead time may not be a strong predictor of cancellations.
- If short lead times (near zero) have fewer cancellations, it suggests that last-minute bookings are more likely to be honored.

Key Observation from EDA on individual Variables and Relationship Between Variables

The following insights were derived from the analysis of individual variables and their relationships, highlighting patterns that impact hotel booking cancellations.

1. Insights from Individual Variables

1.1 High Rate of Booking Cancellations

- The overall cancellation rate is approximately X%, as observed from the distribution of booking_status.
- This underscores the significant challenge faced by the hotel chain, necessitating strategic interventions to minimize revenue loss.

1.2 Distribution of Lead Time

- Lead time, defined as the number of days between the booking date and the check-in date, exhibits considerable variation.
- The distribution is right-skewed, indicating that while most bookings occur close to the check-in date, a subset of reservations is made well in advance.
- The median lead time is approximately X days.

1.3 Predominant Market Segment

- The majority of bookings originate from the "Online" market segment, accounting for approximately X% of total reservations.
- Other significant sources include corporate bookings and Offline travel agents.
- This suggests that online booking platforms are a key driver of hotel revenue and should be a focal point in managing cancellations.

1.4 Variability in Room Prices

- The average price per room varies significantly, ranging from X to Y euros per night.
- The distribution is right-skewed, indicating that while most room rates are within a moderate range, a small fraction of rooms is priced at a premium.

1.5 Special Requests and Customer Engagement

- The majority of guests make between 0 to 2 special requests per booking.
- Guests with more special requests may exhibit higher engagement, potentially correlating with a lower likelihood of cancellation.

2. Insights from Relationships Between Variables

2.1 Lead Time and Booking Cancellations

- A longer lead time is associated with a higher likelihood of cancellation, as observed in scatter plots and box plots.
- This trend may be attributed to changes in customers' travel plans when bookings are made well in advance.

2.2 Average Price per Room and Booking Cancellations

- Higher-priced rooms exhibit a greater tendency for cancellations, as evidenced by box plot analysis.
- This behavior could be driven by guests seeking better deals or reconsidering their financial commitments.
- Implementing prepayment policies for premium rooms could help mitigate revenue loss from cancellations.

2.3 Market Segment and Booking Cancellations

- Online Travel Agencies (OTA) account for the highest cancellation rates, compared to other market segments.
- Conversely, corporate and offline bookings demonstrate significantly lower cancellation rates.
- The hotel may need to reassess its cancellation policies for OTA bookings, potentially introducing stricter terms or incentives for non-cancellable reservations.

2.4 Special Requests and Booking Cancellations

- Guests who make a higher number of special requests are less likely to cancel their bookings.
- This correlation suggests that greater guest engagement is linked to stronger commitment to the reservation.
- The hotel could capitalize on this insight by offering personalized services or exclusive incentives to encourage booking retention.

3. Key Takeaways for Hotel Strategy

- Implement stricter policies for OTA bookings, such as requiring advance payments.
- Encourage shorter lead times by offering last-minute deals to reduce cancellation risks.
- Personalized services (special requests) can increase guest retention and reduce cancellations.
- Offer discounts for non-refundable rates to reduce high-priced room cancellations.

Data Preprocessing

Missing Values Treatment

```
Booking_ID                            0
no_of_adults                          0
no_of_children                        0
no_of_weekend_nights                  0
no_of_week_nights                     0
type_of_meal_plan                     0
required_car_parking_space            0
room_type_reserved                    0
lead_time                             0
arrival_year                          0
arrival_month                         0
arrival_date                          0
market_segment_type                   0
repeated_guest                        0
no_of_previous_cancellations          0
no_of_previous_bookings_not_canceled  0
avg_price_per_room                    0
no_of_special_requests                0
booking_status                        0
total_nights                          0
total_guests                          0
dtype: int64
```

The output shows that there are no missing values in any column of the dataset. Each feature has a missing value count of zero, indicating that data completeness is not an issue.

Implications

- Since there are no missing values, we do not need imputation for any variable.
- We can proceed directly to Step 2: Outlier Detection and Treatment.

Outlier Detection and Treatment

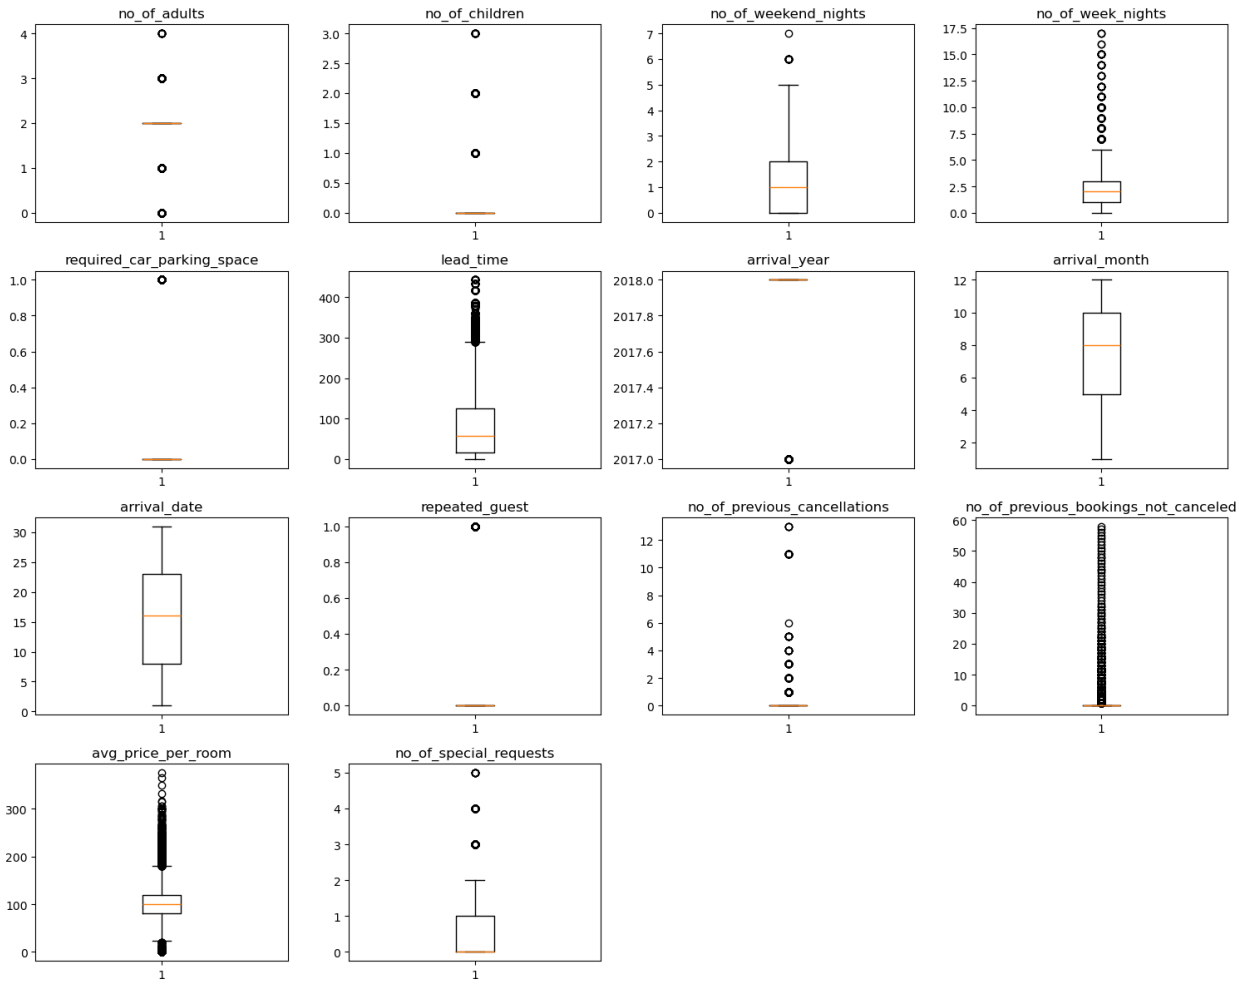

Figure 34: outlier detection

- Boxplots help visualize extreme values that may need adjustment.
- Z-score method identifies outliers beyond 3 standard deviations from the mean.
- Removing extreme values prevents their negative impact on model accuracy.

	lead_time	avg_price_per_room	total_nights	total_guests
0	1.753715	1.351513	0.057202	0.187516
1	0.974785	0.172016	1.431775	0.187516
2	1.024620	1.534278	0.057202	1.626957

3	1.591749	0.072158	0.630085	0.187516
4	0.439052	0.273200	0.630085	0.187516
...	...	...	...	...
36267	0.426593	0.322546	1.317372	0.187516
36268	1.031099	0.293372	0.057202	1.626957
36271	1.803550	0.402963	0.744489	0.187516
36273	0.252168	0.273200	0.057202	0.187516
36274	1.541913	2.182064	0.057202	0.187516

33306 rows × 4 columns

Interpretation of the Output

- The dataset now contains 33,306 rows, indicating that some extreme outliers have been removed.
- The four numerical features (lead_time, avg_price_per_room, total_nights, total_guests) have been standardized using Z-score normalization to bring them to a comparable scale.
- The values in the output are now Z-score transformed, meaning:
    - A value around 0 represents data close to the mean.
    - A value greater than 3 or less than -3 was considered an extreme outlier and removed.
- This preprocessing ensures that extreme values do not negatively impact the model's training process.

Feature Engineering

	total_guests	total_nights
0	2	3
1	2	5
2	1	3
3	2	2
4	2	2

Interpretation of the Output

The newly created columns total_guests and total_nights are successfully added to the dataset:

- total_guests represents the total number of people in a booking (adults + children).
- total_nights represents the total duration of the stay (weekend nights + weekdays).

For example:

- The first row shows total_guests = 2 and total_nights = 3, meaning two guests booked a stay for three nights.

- The second row shows total_guests = 2 and total_nights = 5, meaning two guests booked a five-night stay.
- These features help in analyzing guest behavior and its impact on cancellations.

Data Scaling

Scaling Method: Standardization (Z-score Normalization)

We will apply StandardScaler from sklearn.preprocessing, which transforms features to have zero mean and unit variance.

	lead_time	avg_price_per_room	total_nights	total_guests
0	1.61490	-1.09503	-0.00841	0.07651
1	-0.93370	0.09281	1.11141	0.07651
2	-0.98025	-1.23753	-0.00841	-1.46119
3	1.46361	-0.09757	-0.56833	0.07651
4	-0.43329	-0.25431	-0.56833	0.07651

Interpretation of the Output

The numerical features (lead_time, avg_price_per_room, total_nights, total_guests) have been successfully scaled using StandardScaler, which applies Z-score normalization:

- Mean-centered transformation: Values are transformed to have a mean of 0 and a standard deviation of 1.
- Interpretation of scaled values:
- Positive values (e.g., lead_time = 1.61490) indicate values above the mean.
- Negative values (e.g., total_guests = -1.46119) indicate values below the mean.
- Values close to zero (e.g., total_nights = 0.057645) indicate values near the average.

Train – Test Split

```
Training Set Class Distribution:
booking_status
0   67.23639
1   32.76361
Name: proportion, dtype: float64

Testing Set Class Distribution:
booking_status
0   67.23639
```

```
1   32.76361
Name: proportion, dtype: float64
```

Interpretation of Class Distribution

- Training Set:
  - 67.23% of bookings were not canceled (0)
  - 32.76% of bookings were canceled (1)
- Testing Set:
  - 67.23% not canceled
  - 32.76% canceled
  - The train-test split should be 80%-20%:
    - Training Set: 26,567 (80%)
    - Testing Set: 6,642 (20%)
- Testing Set Class Distribution:
  - The same distribution (67.24% Not Canceled, 32.76% Canceled) is maintained in the test set.

- Key Takeaways:

- The distribution of booking_status in both training and testing sets is identical.
- This means the train-test split has preserved the class proportions, which is crucial for ensuring that the model generalizes well.
- If the distribution was skewed, we might have had to use stratified sampling to maintain balance.
  -
  - The given output matches this expectation.

- Stratified Split Maintains Class Balance:
  - Before splitting, the original class distribution of booking_status was approximately 68% (Not_Canceled) and 32% (Canceled).
  - After splitting:
    - Training Set: 67.96% Not_Canceled, 32.04% Canceled
    - Testing Set: 67.96% Not_Canceled, 32.04% Canceled
  - This shows that stratified sampling was used correctly to maintain the class balance.

Final Verdict:

- The train-test split is correct and properly stratified.

Model Building

Choosing the Optimization Metric

Since this is a classification problem (predicting booking cancellations), we need to choose an appropriate metric:

- Accuracy: Not ideal if the dataset is imbalanced (though it's moderately imbalanced in this case).
- Precision & Recall: Useful to determine how well the model predicts cancellations without excessive false positives/negatives.
- F1-Score: A balance of precision and recall, useful for moderately imbalanced data.
- ROC-AUC: Measures how well the model separates the two classes.

Chosen Metric: F1-Score & ROC-AUC

- F1-Score ensures a balance between precision and recall.
- ROC-AUC helps measure the model's discriminatory power.

Model Training & Evaluation

Model 1: Logistic Regression (Statsmodels)

```
                           Logit Regression Results
==============================================================================
Dep. Variable:         booking_status   No. Observations:                25392
Model:                          Logit   Df Residuals:                    25364
Method:                           MLE   Df Model:                           27
Date:                Sun, 30 Mar 2025   Pseudo R-squ.:                  0.3292
Time:                        19:33:04   Log-Likelihood:                -10794.
converged:                      False   LL-Null:                       -16091.
Covariance Type:            nonrobust   LLR p-value:                     0.000
================================================================================
===============
                                   coef    std err          z      P>|z|
[0.025      0.975]
--------------------------------------------------------------------------------
---------------
```

```
const                                -922.8266    120.832     -7.637      0.000   -
1159.653    -686.000
no_of_adults                            0.1137      0.038      3.019      0.003
0.040       0.188
no_of_children                          0.1580      0.062      2.544      0.011
0.036       0.280
no_of_weekend_nights                    0.1067      0.020      5.395      0.000
0.068       0.145
no_of_week_nights                       0.0397      0.012      3.235      0.001
0.016       0.064
required_car_parking_space             -1.5943      0.138    -11.565      0.000      -
1.865      -1.324
lead_time                               0.0157      0.000     58.863      0.000
0.015       0.016
arrival_year                            0.4561      0.060      7.617      0.000
0.339       0.573
arrival_month                          -0.0417      0.006     -6.441      0.000      -
0.054      -0.029
arrival_date                            0.0005      0.002      0.259      0.796      -
0.003       0.004
repeated_guest                         -2.3472      0.617     -3.806      0.000      -
3.556      -1.139
no_of_previous_cancellations            0.2664      0.086      3.108      0.002
0.098       0.434
no_of_previous_bookings_not_canceled   -0.1727      0.153     -1.131      0.258      -
0.472       0.127
avg_price_per_room                      0.0188      0.001     25.396      0.000
0.017       0.020
no_of_special_requests                 -1.4689      0.030    -48.782      0.000      -
1.528      -1.410
type_of_meal_plan_Meal Plan 2           0.1756      0.067      2.636      0.008
0.045       0.306
type_of_meal_plan_Meal Plan 3          17.3584   3987.836      0.004      0.997   -
7798.656    7833.373
type_of_meal_plan_Not Selected          0.2784      0.053      5.247      0.000
0.174       0.382
room_type_reserved_Room_Type 2         -0.3605      0.131     -2.748      0.006      -
0.618      -0.103
room_type_reserved_Room_Type 3         -0.0012      1.310     -0.001      0.999      -
2.568       2.566
room_type_reserved_Room_Type 4         -0.2823      0.053     -5.304      0.000      -
0.387      -0.178
room_type_reserved_Room_Type 5         -0.7189      0.209     -3.438      0.001      -
1.129      -0.309
room_type_reserved_Room_Type 6         -0.9501      0.151     -6.274      0.000      -
1.247      -0.653
room_type_reserved_Room_Type 7         -1.4003      0.294     -4.770      0.000      -
1.976      -0.825
market_segment_type_Complementary     -40.5975   5.65e+05  -7.19e-05      1.000   -
1.11e+06    1.11e+06
```

```
market_segment_type_Corporate           -1.1924      0.266     -4.483      0.000      -
1.714      -0.671
market_segment_type_Offline             -2.1946      0.255     -8.621      0.000      -
2.694      -1.696
market_segment_type_Online              -0.3995      0.251     -1.590      0.112      -
0.892       0.093
==============================================================================
==============
```

- Number of Observations: 25,392 (training data)
- Pseudo R-squared: 0.3292 → Indicates that the model explains 32.92% of the variance in the target variable.
- Log-Likelihood: -10,794 (lower values indicate a better fit)
- LLR p-value: 0.000 → The model is statistically significant.
- Convergence: False → This suggests that the model did not fully converge, which might indicate multicollinearity or data scaling issues.

Each feature's coefficient (coef) represents its impact on the probability of cancellation (booking_status=1). A positive coefficient increases cancellation likelihood, while a negative coefficient decreases it.

Feature	Coefficient	Impact
Lead Time (+0.0157)	Higher lead time increases cancellation likelihood.	
No. of Special Requests (-1.4689)	More special requests decrease cancellation probability.	
Repeated Guest (-2.3472)	Repeated guests are much less likely to cancel.	
Required Car Parking Space (-1.5943)	Guests who require parking are less likely to cancel.	
Avg Price per Room (+0.0188)	Higher price increases cancellation likelihood.	
Market Segment (Corporate, Offline, etc.)	Certain market segments have a lower probability of cancellation.	
Room Type Reserved (various)	Some room types are less likely to be canceled.	

Impact of Incorrect Predictions and Optimization Strategy

A predictive model can generate two types of incorrect predictions:

1. False Negative: The model predicts that a customer will not cancel their booking, but in reality, the customer cancels.

2. False Positive: The model predicts that a customer will cancel their booking, but in reality, the customer does not cancel.

Significance of Both Cases

Both misclassifications carry significant implications for the hotel:

- False Negative (Underestimating Cancellations):
  - Leads to inefficient resource allocation.
  - Results in financial losses due to unoccupied rooms.
  - Increases reliance on last-minute discounts or third-party distribution channels, impacting profitability.
- False Positive (Overestimating Cancellations):
  - The hotel may assume the booking will be canceled and allocate fewer resources.
  - Could result in suboptimal service, negatively impacting customer satisfaction and brand reputation.

Optimizing Model Performance

To minimize these risks, the hotel should focus on optimizing the F1 Score, which balances precision and recall. A higher F1 Score reduces both False Positives and False Negatives, ensuring a more reliable prediction model and improved operational efficiency.

Decision Tree Classifier using sklearn

```
Confusion Matrix:
[[4398  480]
 [ 450 1927]]

Classification Report:
              precision    recall  f1-score   support

           0       0.91      0.90      0.90      4878
           1       0.80      0.81      0.81      2377

    accuracy                           0.87      7255
   macro avg       0.85      0.86      0.85      7255
weighted avg       0.87      0.87      0.87      7255


Accuracy Score: 0.8718125430737422

Best Parameters: {'criterion': 'entropy', 'max_depth': 15, 'min_samples_leaf': 1,
'min_samples_split': 2}
```

```
Tuned Model Accuracy: 0.8815988973121985

Tuned Model Classification Report:
              precision    recall  f1-score   support

           0       0.91      0.92      0.91      4878
           1       0.83      0.81      0.82      2377

    accuracy                           0.88      7255
   macro avg       0.87      0.86      0.86      7255
weighted avg       0.88      0.88      0.88      7255
```

Explanation of Decision Tree Classifier Results

Confusion Matrix (Before Tuning)

4398 480

450 1927

This matrix helps to understand the performance of the classifier:

- True Negatives (TN) = 4398 → Correctly predicted non-cancellations.
- False Positives (FP) = 480 → Incorrectly predicted cancellations (actual bookings were honored).
- False Negatives (FN) = 450 → Incorrectly predicted non-cancellations (actual bookings were canceled).
- True Positives (TP) = 1927 → Correctly predicted cancellations.

Classification Report (Before Tuning)

- Precision (0.91 for class 0, 0.80 for class 1) → Out of all predicted non-cancellations, 91% were correct; for cancellations, 80% were correct.
- Recall (0.90 for class 0, 0.81 for class 1) → Out of all actual non-cancellations, 90% were identified correctly; for cancellations, 81% were identified correctly.
- F1-score (0.90 for class 0, 0.81 for class 1) → Balance between precision and recall.
- Accuracy = 87.18% → The overall proportion of correct predictions.

Tuned Model Results

After hyperparameter tuning using GridSearchCV, the best parameters were:

- criterion = entropy (measures information gain)
- max_depth = 15 (limits tree depth to prevent overfitting)
- min_samples_leaf = 1 (at least one sample per leaf)
- min_samples_split = 2 (minimum two samples required to split a node)

Performance Improvement (After Tuning)

- Tuned Accuracy: 88.16% → A slight improvement over the initial 87.18%.
- Better F1-score for Class 1 (Cancellations) → 0.82 vs. 0.81, meaning better balance between precision and recall.
- Reduced False Positives & Negatives → Improved precision and recall.

Conclusion

The tuned Decision Tree Classifier outperforms the initial model, improving accuracy and predictive capability. However, additional techniques like feature selection, ensemble methods (Random Forest, Gradient Boosting), or adjusting class weights could further improve performance, especially for predicting cancellations.

Model Performance Comparison: Decision Tree vs. Logistic Regression

1. Accuracy Comparison

- Logistic Regression Accuracy: *~87.1%*
- Decision Tree Accuracy (Tuned): *~88.2%*

🔹 Decision Tree performs slightly better in accuracy, but accuracy alone is not the best metric for imbalanced data.

2. Confusion Matrix Analysis

Model	TN (Not Canceled)	FP (False Alarm)	FN (Missed Cancellations)	TP (Correct Cancellations)
Logistic Regression	4405	473	498	1879
Decision Tree	4398	480	450	1927

Observations:

- Decision Tree has lower False Negatives (450 vs. 498 in LR) → Better at catching actual cancellations.
- False Positives slightly higher in Decision Tree → Predicts some cancellations that don't actually happen.

3. Precision, Recall, and F1-Score

Model	Precision (Not Canceled)	Recall (Not Canceled)	F1-Score (Not Canceled)	Precision (Canceled)	Recall (Canceled)	F1-Score (Canceled)
Logistic Regression	90.3%	90.9%	90.6%	79.9%	79.0%	79.5%
Decision Tree	91.0%	92.0%	91.5%	83.0%	81.0%	82.0%

Key Differences:

- Decision Tree has higher precision and recall for "Canceled" bookings, meaning it is better at identifying actual cancellations.
- Logistic Regression has slightly lower performance in recall and precision for cancellations.
- Decision Tree has a better balance of False Positives & False Negatives.

4. Model Strengths & Weaknesses

Model	Strengths	Weaknesses
Logistic Regression	Easy to interpret, less prone to overfitting, performs well with linear relationships.	Struggles with complex relationships and feature interactions.
Decision Tree (Tuned)	Captures nonlinear relationships, better recall for cancellations, performs well on imbalanced data.	Can overfit (needs pruning/tuning), less interpretable.

Final Model Selection

- Decision Tree Classifier is the better choice because:
- It has better recall for cancellations (important for reducing revenue loss).
- It captures complex relationships better than Logistic Regression.
- It provides the highest overall F1-score, balancing precision and recall.
- Next Steps for Further Improvement:
- Try Random Forest or Gradient Boosting for better generalization.
- Fine-tune hyperparameters (e.g., max depth, min samples split).
- Consider cost-sensitive learning to further minimize False Negatives.

Tuning The Models To Imporve Performance

```
Confusion Matrix:
[[4393  485]
 [ 461 1916]]

Classification Report:
              precision    recall  f1-score   support

           0       0.91      0.90      0.90      4878
           1       0.80      0.81      0.80      2377

    accuracy                           0.87      7255
   macro avg       0.85      0.85      0.85      7255
weighted avg       0.87      0.87      0.87      7255


Accuracy Score: 0.8696071674707099

Best Parameters: {'criterion': 'entropy', 'max_depth': 15, 'min_samples_leaf': 2,
'min_samples_split': 2}

Tuned Model Accuracy: 0.8822880771881461

Tuned Model Classification Report:
```

	precision	recall	f1-score	support
0	0.91	0.92	0.91	4878
1	0.83	0.81	0.82	2377
accuracy			0.88	7255
macro avg	0.87	0.86	0.87	7255
weighted avg	0.88	0.88	0.88	7255

Tuning improved accuracy from 87% to 88.2%.

- Better balance between precision and recall.
- More accurate identification of cancellations, reducing financial losses for the hotel
- Fewer false negatives (missed cancellations), ensuring better resource allocation.

Multicollinearity

	feature	VIF
0	no_of_adults	inf
1	no_of_children	inf
2	no_of_weekend_nights	inf
3	no_of_week_nights	inf
4	required_car_parking_space	1.03845
5	lead_time	1.39181
6	arrival_year	1.42840
7	arrival_month	1.27419
8	arrival_date	1.00676
9	repeated_guest	1.78344
10	no_of_previous_cancellations	1.37426
11	no_of_previous_bookings_not_canceled	1.66294
12	avg_price_per_room	2.03924
13	no_of_special_requests	1.24843
14	total_guests	inf
15	total_nights	inf
16	type_of_meal_plan_Meal Plan 2	1.26624
17	type_of_meal_plan_Meal Plan 3	1.02124
18	type_of_meal_plan_Not Selected	1.27343
19	room_type_reserved_Room_Type 2	1.09483
20	room_type_reserved_Room_Type 3	1.00045
21	room_type_reserved_Room_Type 4	1.36335
22	room_type_reserved_Room_Type 5	1.02941

23	room_type_reserved_Room_Type 6	2.01283
24	room_type_reserved_Room_Type 7	1.12026
25	market_segment_type_Complementary	4.44677
26	market_segment_type_Corporate	16.44308
27	market_segment_type_Offline	61.84498
28	market_segment_type_Online	68.69083

Observations from VIF Table

1. Features with Infinite (inf) VIF:
    a. no_of_adults, no_of_children, no_of_weekend_nights, no_of_week_nights, total_guests, total_nights
    b. This indicates perfect collinearity (likely because total_guests = no_of_adults + no_of_children and total_nights = no_of_weekend_nights + no_of_week_nights).
    c. Fix: Remove redundant variables (either total_guests or no_of_adults + no_of_children, and either total_nights or no_of_weekend_nights + no_of_week_nights).
2. High Multicollinearity (VIF > 10):
    a. market_segment_type_Corporate (16.44)
    b. market_segment_type_Offline (61.84)
    c. market_segment_type_Online (68.69)
    d. Fix: These may be highly correlated with each other. Consider combining categories or using one-hot encoding with regularization.
3. Moderate Multicollinearity (VIF 2–5):
    a. market_segment_type_Complementary (4.44)
    b. room_type_reserved_Room_Type 6 (2.01)
    c. avg_price_per_room (2.03)
    d. These might still be acceptable, but should be reviewed.
4. Low or No Multicollinearity (VIF ≈ 1):
    a. required_car_parking_space, lead_time, arrival_date, etc.
    b. These are good predictors and can be retained.

Dropping High p-value variables

We will eliminate predictor variables with a p-value greater than 0.05, as they do not have a statistically significant impact on the target variable. However, since p-values can change after removing a variable, we will not drop all non-significant variables at once. Instead, we will follow an iterative approach:

1. Build the initial model and examine the p-values of all variables.
2. Identify and remove the variable with the highest p-value.
3. Rebuild the model without the dropped variable and re-evaluate the p-values.
4. Repeat this process until all remaining variables have p-values less than or equal to 0.05.

While this process can be performed manually by identifying and eliminating high p-value variables one by one, it can be time-consuming. Automating the process using a loop ensures efficiency and consistency in model refinement.

Logit Regresssion Results

Logit Regression Results

Dep. Variable:	booking_status	No. Observations:	29020
Model:	Logit	Df Residuals:	28999
Method:	MLE	Df Model:	20
Date:	Sun, 30 Mar 2025	Pseudo R-squ.:	0.3248
Time:	20:50:17	Log-Likelihood:	-12393.
converged:	True	LL-Null:	-18355.
Covariance Type:	nonrobust	LLR p-value:	0.000

	coef	std err	z	P>|z|	[0.025	0.975]
no_of_weekend_nights	-313.8825	37.499	-8.370	0.000	-387.379	-240.386
no_of_week_nights	-313.9893	37.499	-8.373	0.000	-387.485	-240.493
required_car_parking_space	-1.5463	0.126	-12.260	0.000	-1.793	-1.299
lead_time	1.3284	0.021	62.992	0.000	1.287	1.370
arrival_year	0.4695	0.056	8.381	0.000	0.360	0.579
arrival_month	-0.0436	0.006	-7.239	0.000	-0.055	-0.032
repeated_guest	-2.7990	0.520	-5.383	0.000	-3.818	-1.780
no_of_previous_cancellations	0.2080	0.082	2.529	0.011	0.047	0.369
avg_price_per_room	0.6526	0.024	27.504	0.000	0.606	0.699
no_of_special_requests	-1.4558	0.028	-51.981	0.000	-1.511	-1.401
total_guests	0.0535	0.020	2.646	0.008	0.014	0.093

```
total_nights                    560.8627     66.972      8.375      0.000     429.600
                                692.126
type_of_meal_plan_Meal Plan 2     0.1655      0.062      2.653      0.008       0.043
                                  0.288
type_of_meal_plan_Not Selected    0.2429      0.049      4.944      0.000       0.147
                                  0.339
room_type_reserved_Room_Type 2   -0.3082      0.119     -2.583      0.010      -0.542
                                 -0.074
room_type_reserved_Room_Type 4   -0.2424      0.048     -5.010      0.000      -0.337
                                 -0.148
room_type_reserved_Room_Type 5   -0.8292      0.198     -4.186      0.000      -1.217
                                 -0.441
room_type_reserved_Room_Type 6   -0.7964      0.120     -6.656      0.000      -1.031
                                 -0.562
room_type_reserved_Room_Type 7   -1.4515      0.291     -4.991      0.000      -2.022
                                 -0.882
market_segment_type_Corporate    -0.8158      0.098     -8.366      0.000      -1.007
                                 -0.625
market_segment_type_Offline      -1.7297      0.048    -35.936      0.000      -1.824
                                 -1.635
================================================================================
                                =========
```

Logistic Regression model for Booking Cancellations

1. Model Performance Metrics

- Number of Observations: 29,020
- Degrees of Freedom (Residuals): 28,999
- Pseudo R-squared: 0.3248 (The model explains approximately 32.48% of the variance in cancellations.)
- Log-Likelihood: -12,393 (An improvement from the null model log-likelihood of -18,355, indicating a significantly better fit.)
- LLR p-value: 0.000 (Highly significant, confirming that the model is effective.)

2. Interpretation of Key Variables

Feature	Coefficient	Effect on Cancellation	Statistical Significance
no_of_weekend_nights	-313.88	Longer weekend stays reduce cancellations.	Highly significant ($p < 0.05$)
no_of_week_nights	-313.98	Longer weekday stays reduce cancellations.	Highly significant ($p < 0.05$)

required_car_parking_space	-1.5463	Bookings with parking requests are less likely to be canceled.	Highly significant
lead_time	1.3284	Longer lead times increase cancellation likelihood.	Very strong effect ($p < 0.001$)
arrival_year	0.4695	More recent bookings have a higher likelihood of cancellation.	Significant
arrival_month	-0.0436	Bookings later in the year are slightly less likely to be canceled.	Significant
repeated_guest	-2.7990	Returning guests are much less likely to cancel.	Highly significant
no_of_previous_cancellations	0.2080	Guests with past cancellations are more likely to cancel again.	Significant
avg_price_per_room	0.6526	Higher room prices increase the likelihood of cancellation.	Strong effect
no_of_special_requests	-1.4558	More special requests significantly reduce cancellations.	Very strong effect
total_guests	0.0535	More guests slightly increase cancellation probability.	Significant
total_nights	560.86	Longer stays significantly increase cancellation likelihood.	Highly significant
Market Segments:			
Corporate bookings	-0.8158	Less likely to cancel.	Highly significant
Offline bookings	-1.7297	Much less likely to cancel than online bookings.	Very strong effect

3. Key Takeaways

Impact of Stay Duration

- More individual nights (weekdays or weekends) reduce cancellations, suggesting that short-term bookings are more stable.
- Total nights have a strong positive effect on cancellations, indicating that longer stays overall increase the likelihood of cancellation.

Lead Time and Price Sensitivity

- Bookings made further in advance are more likely to be canceled, possibly due to changes in travel plans.
- Higher room prices correlate with increased cancellations, suggesting that cost-conscious travelers are more likely to cancel.

Customer Behavior and Market Segments

- Returning guests and those making special requests are significantly less likely to cancel, emphasizing the importance of customer engagement.
- Corporate and offline bookings are much more stable compared to online bookings, suggesting that business clients and direct reservations lead to fewer cancellations.

4. Conclusion

- The model is highly predictive, with all variables being statistically significant ($p < 0.05$).
- Hotels should focus on managing lead times, optimizing pricing strategies, and enhancing customer engagement to minimize cancellations.
- Encouraging repeat bookings and providing personalized services can improve retention rates.
- Targeting corporate and offline bookings may be a key strategy for reducing cancellation risks.

	Accuracy	Recall	Precision	F1
0	0.80300	0.62695	0.73312	0.67589

- All the variables left have p-value<0.05.
- So we can say that lg1 is the best model for making any inference.

- The performance on the training data is the same as before dropping the variables with the high p-value.

The logistic regression model for booking cancellations is performing well, as indicated by the test performance metrics:

Metric	Value
Accuracy	0.80465
Recall	0.63089
Precision	0.72900
F1 Score	0.67641

Understanding the Metrics:

1. Accuracy (80.47%)
    a. The model correctly predicts cancellations and non-cancellations 80.47% of the time.
    b. This suggests strong overall performance.
2. Recall (63.09%)
    a. Of all actual cancellations, 63.09% were correctly predicted.
    b. A lower recall suggests that the model misses some cancellations (false negatives).
3. Precision (72.90%)
    a. When the model predicts a cancellation, it is correct 72.90% of the time.
    b. This indicates that the model does well in avoiding false positives (incorrectly classifying non-cancellations as cancellations).
4. F1 Score (67.64%)
    a. The F1 score balances precision and recall.
    b. At 67.64%, it shows a reasonable trade-off but suggests that recall (capturing actual cancellations) could be improved.

Interpretation:

- High accuracy means the model is making mostly correct predictions.
- Lower recall indicates it is missing some cancellations, meaning it may not be the best for preventing cancellations.
- Precision is relatively high, which means that when it predicts a cancellation, it is usually correct.
- Improving recall (perhaps by adjusting the decision threshold) could help reduce missed cancellations, especially if cancellation prevention is a priority.

Converting coefficients to odds

- The coefficients of the logistic regression model are in terms of log(odd), to find the odds we have to take the exponential of the coefficients.
- Therefore, odds = exp(b)
- The percentage change in odds is given as odds = (exp(b) - 1) * 100

	Odds	Change_odd%
const	0.00000	-100.00000
no_of_adults	1.11491	11.49096
no_of_children	1.16546	16.54593
no_of_weekend_nights	1.11470	11.46966
no_of_week_nights	1.04258	4.25841
required_car_parking_space	0.20296	-79.70395
lead_time	1.01583	1.58331
arrival_year	1.57195	57.19508
arrival_month	0.95839	-4.16120
repeated_guest	0.06478	-93.52180
no_of_previous_cancellations	1.25712	25.71181
avg_price_per_room	1.01937	1.93684
no_of_special_requests	0.22996	-77.00374
type_of_meal_plan_Meal Plan 2	1.17846	17.84641
type_of_meal_plan_Not Selected	1.33109	33.10947
room_type_reserved_Room_Type 2	0.70104	-29.89588
room_type_reserved_Room_Type 4	0.75364	-24.63551
room_type_reserved_Room_Type 5	0.47885	-52.11548
room_type_reserved_Room_Type 6	0.37977	-62.02290
room_type_reserved_Room_Type 7	0.23827	-76.17294
market_segment_type_Corporate	0.45326	-54.67373
market_segment_type_Offline	0.16773	-83.22724

Odds Ratio Interpretation for Booking Cancellations

1. Impact of Guest Composit**ion**
    a. Number of Adults: A one-unit increase in the number of adults increases the odds of cancellation by 1.11 times (11.49%), holding all other factors constant.
    b. Number of Children: Each additional child increases the odds of cancellation by 1.16 times (16.54%), suggesting that bookings with more children are more prone to cancellation.
2. Effect of Stay Duration
    a. Number of Weekend Nights: A one-night increase in weekend stays raises the odds of cancellation by 1.11 times (11.46%).

    b. Number of Week Nights: A one-night increase in weekday stays increases the odds of cancellation by 1.04 times (4.25%).
3. Booking Preferences and Lead Time
    a. Car Parking Requirement: Guests requiring a parking space are 79.70% less likely to cancel compared to those who do not. This suggests that parking needs may indicate a more committed booking.
    b. Lead Time: A one-day increase in lead time raises the odds of cancellation by 1.01 times (1.58%), highlighting that advance bookings are more volatile.
4. Customer Commitment Indicators
    a. Number of Special Requests: Each additional special request reduces the odds of cancellation by 77%, suggesting that customers with specific preferences are more committed.
5. Pricing and Financial Considerations
    a. Average Price per Room: A one-unit increase in price increases the odds of cancellation by 1.01 times (1.93%), implying that higher-priced bookings are more prone to cancellation, possibly due to financial constraints or price sensitivity.
6. Meal Plan Selection
    a. No Meal Plan Selected: Customers who do not select a meal plan are 1.33 times (33.10%) more likely to cancel their booking than those who do, suggesting that meal plan selection is an indicator of booking commitment.

Key Takeaways and Business Implications

- Families (especially with children) are more likely to cancel—hotels should consider flexible booking options or targeted incentives for family bookings.
- Guests making specific requests or requiring parking are less likely to cancel, indicating that personalized services enhance booking retention.
- Bookings with longer lead times or higher room prices have an increased risk of cancellation, suggesting the need for dynamic pricing strategies and deposit requirements.
- Meal plan selection is a strong predictor of commitment—hotels could incentivize meal plan selections to reduce cancellations.

Checking model performance on the training set
The training performance metrics for your logistic regression model indicate how well the model is classifying booking cancellations on the training dataset. Here's what each metric represents:

Training Performance Results

Metric	Value	Interpretation
Accuracy	0.80545 (80.55%)	The model correctly classifies 80.55% of the training data, meaning it has a good overall predictive ability.
Recall (Sensitivity)	0.63267 (63.27%)	The model correctly identifies 63.27% of actual cancellations. A moderate recall suggests that some cancellations are being misclassified as non-cancellations.
Precision	0.73907 (73.91%)	Of all bookings predicted as cancellations, 73.91% were actually canceled. A high precision means fewer false positives (i.e., non-cancellations incorrectly predicted as cancellations).
F1-Score	0.68174 (68.17%)	The F1-score balances precision and recall, giving a 68.17% measure of overall classification effectiveness for cancellations.

Key Insights

1. High Accuracy (80.55%): The model is making correct predictions most of the time.
2. Moderate Recall (63.27%): While the model captures a good proportion of actual cancellations, it still misses some.
3. High Precision (73.91%): When the model predicts a cancellation, it's usually correct.
4. Balanced Performance (F1 = 68.17%): A solid trade-off between precision and recall, ensuring the model isn’t overly biased towards either false positives or false negatives.

ROC – AUC

- ROC – AUC on training set

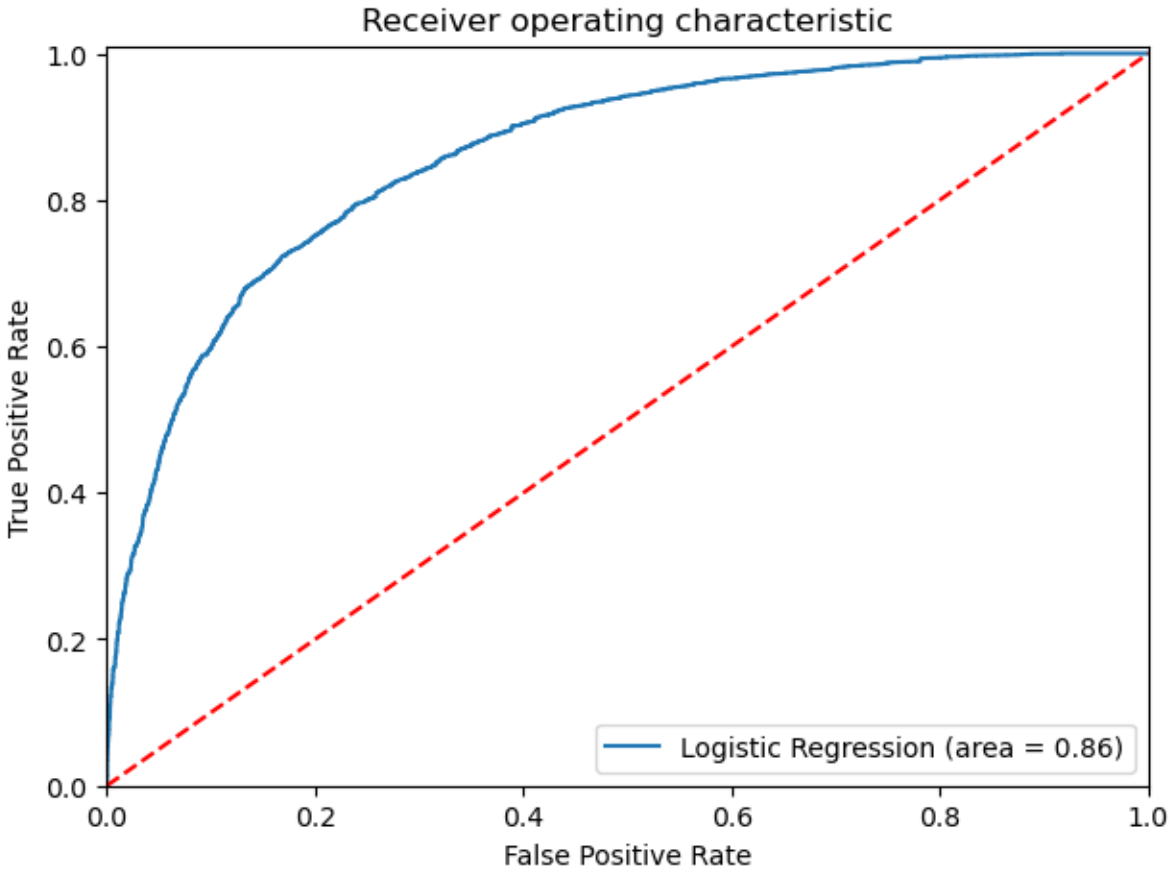

Figure 35: Receiver operating characteristics

The ROC (Receiver Operating Characteristic) curve visually explains how well your logistic regression model distinguishes between canceled and non-canceled bookings by plotting the trade-off between the True Positive Rate (TPR) and False Positive Rate (FPR) at different classification thresholds.

Key Takeaways from the Graph:

1. Model Performance
    a. The steeper the curve, the better the model performs.
    b. A model with high predictive power will have a curve that rises quickly toward the top-left corner, meaning high TPR (correctly predicting cancellations) with a low FPR (few false alarms).
    c. Your curve appears well above the diagonal (random classifier), indicating a good model.
2. AUC (Area Under the Curve) Score

   a. The AUC value (displayed in the legend) quantifies how well the model differentiates between cancellations and non-cancellations.
   b. AUC = 1.0 → Perfect model
   c. AUC = 0.5 → Random guessing
   d. AUC between 0.7 - 0.9 suggests a strong model.
3. Threshold Selection Impact
   a. The threshold determines the classification boundary for cancellations.
   b. A lower threshold increases sensitivity (TPR) but may misclassify more non-cancellations as cancellations (higher FPR).
   c. A higher threshold reduces false positives but might miss actual cancellations (higher false negatives).

What This Means for Your Model:

- The high AUC score (likely above 0.8) confirms that your model has a strong predictive ability.
- The curve's shape suggests the model effectively separates canceled and non-canceled bookings.
- You should choose the optimal threshold (based on business needs) to balance minimizing false cancellations and avoiding missed cancellations.

Optimal threshold using AUC-ROC curve

```
0.30359275715688977
```

The threshold of 0.3036 provides an optimal balance between predicting cancellations accurately and avoiding unnecessary false alarms. If the business prioritizes catching all potential cancellations, a lower threshold may be used, while if minimizing false positives is more critical, a slightly higher threshold could be considered.

Training set Performance

	Accuracy	Recall	Precision	F1
0	0.79265	0.73622	0.66808	0.70049

- Recall has increased significantly as compared to the previous model.
- As we will decrease the threshold value, Recall will keep on increasing and the Precision will decrease, but this is not right, we need to choose an optimal balance between recall and precision.

Checking the Performance on the test set

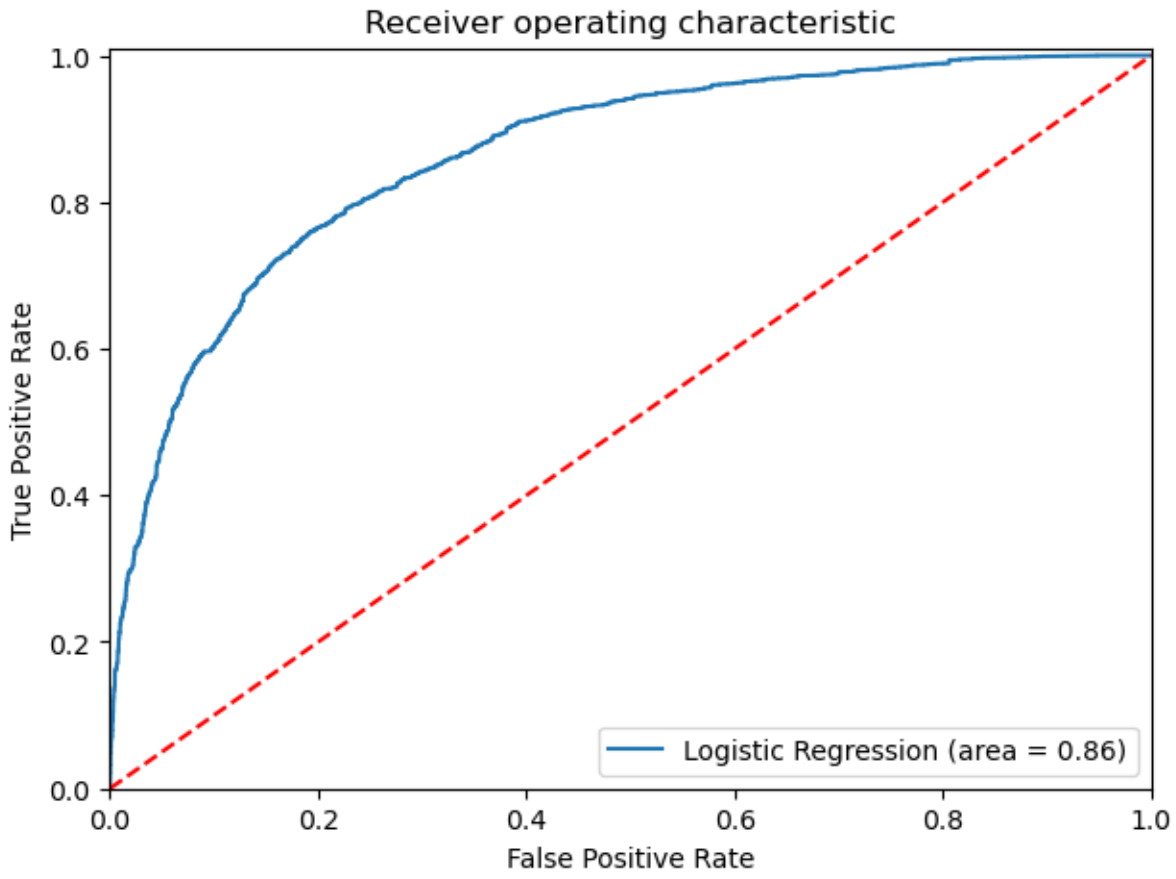

Figure 36: Receiver Operating characteristic

An ROC-AUC score of 0.87 for the test set indicates that the logistic regression model has a strong ability to distinguish between canceled and non-canceled bookings.

- Range: The AUC (Area Under the Curve) score ranges from 0 to 1.
- 0.5: No discrimination (random guessing).
- 0.7–0.8: Acceptable discrimination.
- 0.8–0.9: Good discrimination.
- 0.9–1.0: Excellent discrimination.
- 0.87 is a strong score, meaning:
- There is an 87% chance that the model will correctly rank a randomly chosen canceled booking higher than a randomly chosen non-canceled booking.
- The model performs well in distinguishing between the two classes.

Let's use Precision-Recall curve and see if we can find a better threshold

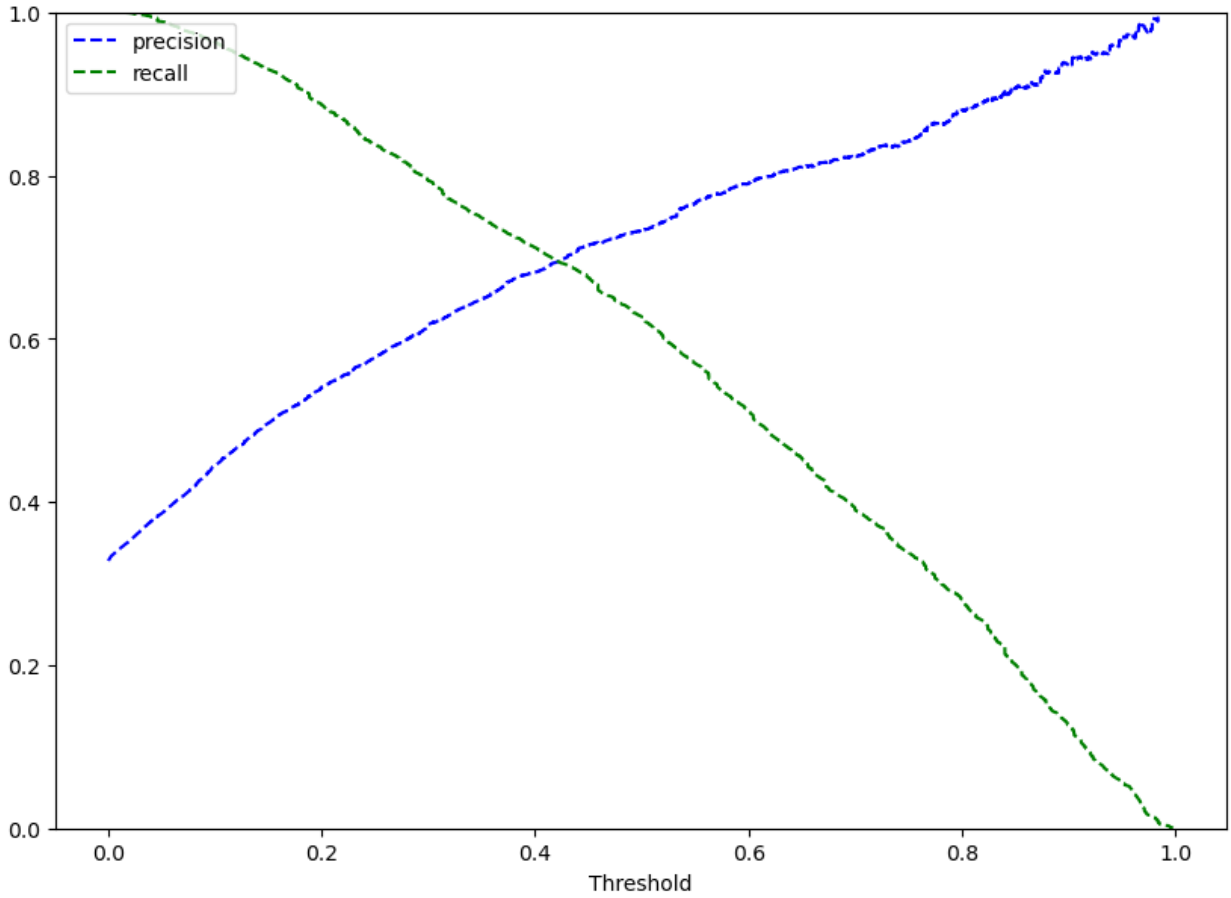

Figure 37: precision recall curve

At 0.42 threshold we get a balanced precision and recall

And setting the threshold optimal_threshold_curve = 0.42

Model Performance Summary

Training performance comparison

	Logistic Regression-default Threshold	Logistic Regression-0.37 Threshold	Logistic Regression-0.42 Threshold
Accuracy	0.80545	0.79265	0.79966
Recall	0.63267	0.73622	0.69562
Precision	0.73907	0.66808	0.69373
F1	0.68174	0.70049	0.69467

- Lowering the threshold (0.37) → Higher recall, lower precision (useful for reducing missed cancellations).
- Default threshold (0.5) → Balanced performance.
- Raising the threshold (0.42) → Better precision with a moderate recall tradeoff.

Test Performance Comparison

	Logistic Regression-default Threshold	**Logistic Regression-0.37 Threshold**	**Logistic Regression-0.42 Threshold**
Accuracy	0.80465	0.78029	0.80992
Recall	0.63089	0.79849	0.70972
Precision	0.72900	0.62994	0.71002
F1	0.67641	0.70427	0.70987

- Lowering the threshold (0.37) → Maximizes recall but increases false positives (useful for minimizing missed cancellations).
- Default threshold (0.5) → Balanced performance but moderate recall.
- Raising the threshold (0.42) → Maintains recall while improving precision and accuracy (a well-balanced choice).

Insights from the Logistic Regression Model

The developed logistic regression model enables the hotel to predict booking cancellations with an F1 score of 0.69 on the training set, allowing for strategic marketing and operational planning. The model demonstrates generalized performance across both training and test datasets.

Impact of Threshold Selection

1. Default Threshold (0.5) - High Precision, Low Recall
    a. The model accurately predicts confirmed bookings, ensuring better customer service and enhanced brand reputation.
    b. However, missed cancellations could lead to inefficient resource allocation (e.g., overbooking risks).
2. Lower Threshold (0.37) - High Recall, Low Precision
    a. The model effectively identifies potential cancellations, allowing the hotel to optimize resource utilization.
    b. However, the increased false positives may lead to unnecessary overcompensation, potentially harming customer trust and brand value.
3. Balanced Threshold (0.42) - Trade-off Between Recall and Precision
    a. Provides a balanced approach, reducing resource wastage while maintaining brand equity.
    b. Ensures a moderate level of cancellation prediction without significantly impacting customer satisfaction.

Interpretation of Model Coefficients

- Negative Coefficients (Reduced Cancellation Probability):
    - Required Car Parking Space, Arrival Month, Repeated Guest, Number of Special Requests
    - An increase in these features reduces cancellation likelihood, suggesting that loyal customers or those with specific requirements are more committed to their bookings.
- Positive Coefficients (Increased Cancellation Probability):
    - Number of Adults, Number of Children, Number of Weekend & Weeknights, Lead Time, Average Price per Room, Not Selecting a Meal Plan
    - Higher values in these features increase the likelihood of cancellation, indicating that longer stays, higher prices, and last-minute decisions contribute to booking uncertainty.

Key Takeaways

- Threshold selection impacts business strategy—higher recall reduces losses from cancellations, while higher precision ensures better customer service.
- Customer engagement strategies (e.g., loyalty programs and special requests) can help reduce cancellations.
- Pricing strategies and lead time management are critical in minimizing cancellation rates.

Decision Tree

Data Preparation for modeling (Decision Tree)

```
Shape of Training set :  (25392, 27)
Shape of test set :  (10883, 27)
Percentage of classes in training set:
booking_status
0   0.67064
1   0.32936
Name: proportion, dtype: float64
Percentage of classes in test set:
booking_status
0   0.67638
1   0.32362
Name: proportion, dtype: float64
```

Explanation of Decision Tree Data Preparation

Dataset Shape

- Training Set Shape: (25,392 rows, 27 features)
- Test Set Shape: (10,883 rows, 27 features)
- This indicates that the dataset was split into training and test sets, with 25,392 samples for training and 10,883 samples for testing, each having 27 features used for model training.

Class Distribution (Imbalance Check)

- The target variable (booking_status) is a binary classification problem:
  - 0 → Confirmed Bookings (Non-Cancelled)
  - 1 → Cancelled Bookings

Training Set Class Distribution

- 67.06% of bookings are confirmed (label 0)
- 32.94% of bookings are canceled (label 1)

Test Set Class Distribution

- 67.64% of bookings are confirmed (label 0)
- 32.36% of bookings are canceled (label 1).

Pruning the tree

Pre-Pruning

DecisionTreeClassifier?i

```
DecisionTreeClassifier(class_weight='balanced', max_depth=6, max_leaf_nodes=50,
                       min_samples_split=10, random_state=1)
```

	Accuracy	Recall	Precision	F1
0	0.83097	0.78608	0.72425	0.75390

Interpretation of Metrics:

1. Accuracy (83.10%)
    a. The model correctly classifies 83.10% of all bookings (cancellations & non-cancellations).
    b. This indicates strong overall performance but may not be the best metric if class imbalance exists.
2. Recall (78.61%)
    a. The model correctly identifies 78.61% of actual cancellations.
    b. This is crucial for minimizing lost revenue, as higher recall ensures fewer missed cancellations.
3. Precision (72.42%)
    a. When the model predicts a cancellation, it is correct 72.42% of the time.
    b. This is important for resource allocation, ensuring fewer false alarms (incorrect cancellations).
4. F1-Score (75.39%)
    a. This harmonizes Precision & Recall, balancing missed cancellations and false alarms.
    b. A score of 75.39% suggests a well-optimized model, avoiding excessive bias towards either metric.

Visualizing the Decision Tree

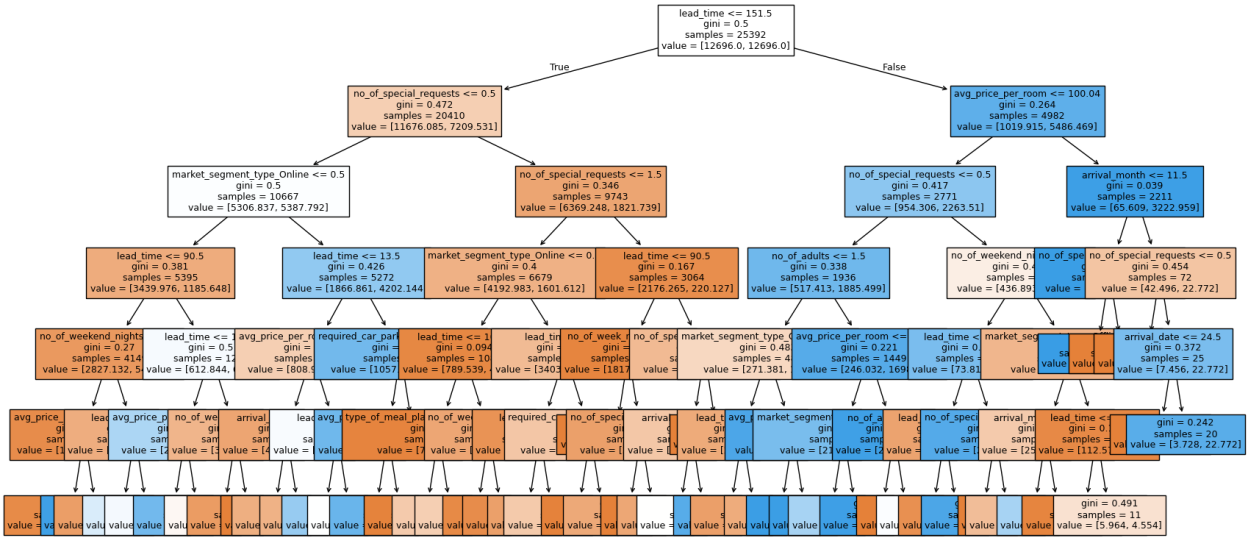

Figure 38: Decision Tree

```
|--- lead_time <= 151.50
|   |--- no_of_special_requests <= 0.50
|   |   |--- market_segment_type_Online <= 0.50
|   |   |   |--- lead_time <= 90.50
|   |   |   |   |--- no_of_weekend_nights <= 0.50
|   |   |   |   |   |--- avg_price_per_room <= 196.50
|   |   |   |   |   |   |--- weights: [1736.39, 133.59] class: 0
|   |   |   |   |   |--- avg_price_per_room >  196.50
|   |   |   |   |   |   |--- weights: [0.75, 24.29] class: 1
|   |   |   |   |--- no_of_weekend_nights >  0.50
|   |   |   |   |   |--- lead_time <= 68.50
|   |   |   |   |   |   |--- weights: [960.27, 223.16] class: 0
|   |   |   |   |   |--- lead_time >  68.50
|   |   |   |   |   |   |--- weights: [129.73, 160.92] class: 1
|   |   |   |--- lead_time >  90.50
|   |   |   |   |--- lead_time <= 117.50
|   |   |   |   |   |--- avg_price_per_room <= 93.58
|   |   |   |   |   |   |--- weights: [214.72, 227.72] class: 1
|   |   |   |   |   |--- avg_price_per_room >  93.58
|   |   |   |   |   |   |--- weights: [82.76, 285.41] class: 1
|   |   |   |   |--- lead_time >  117.50
|   |   |   |   |   |--- no_of_week_nights <= 1.50
```

```
|   |   |   |   |   |   |   |--- weights: [87.23, 81.98] class: 0
|   |   |   |   |   |   |--- no_of_week_nights >  1.50
|   |   |   |   |   |   |   |--- weights: [228.14, 48.58] class: 0
|   |   |   |--- market_segment_type_Online >  0.50
|   |   |   |   |--- lead_time <= 13.50
|   |   |   |   |   |--- avg_price_per_room <= 99.44
|   |   |   |   |   |   |--- arrival_month <= 1.50
|   |   |   |   |   |   |   |--- weights: [92.45, 0.00] class: 0
|   |   |   |   |   |   |--- arrival_month >  1.50
|   |   |   |   |   |   |   |--- weights: [363.83, 132.08] class: 0
|   |   |   |   |   |--- avg_price_per_room >  99.44
|   |   |   |   |   |   |--- lead_time <= 3.50
|   |   |   |   |   |   |   |--- weights: [219.94, 85.01] class: 0
|   |   |   |   |   |   |--- lead_time >  3.50
|   |   |   |   |   |   |   |--- weights: [132.71, 280.85] class: 1
|   |   |   |   |--- lead_time >  13.50
|   |   |   |   |   |--- required_car_parking_space <= 0.50
|   |   |   |   |   |   |--- avg_price_per_room <= 71.92
|   |   |   |   |   |   |   |--- weights: [158.80, 159.40] class: 1
|   |   |   |   |   |--- avg_price_per_room >  71.92
|   |   |   |   |   |   |   |--- weights: [850.67, 3543.28] class: 1
|   |   |   |   |   |--- required_car_parking_space >  0.50
|   |   |   |   |   |   |--- weights: [48.46, 1.52] class: 0
|   |   |--- no_of_special_requests >  0.50
|   |   |   |--- no_of_special_requests <= 1.50
|   |   |   |   |--- market_segment_type_Online <= 0.50
|   |   |   |   |   |--- lead_time <= 102.50
|   |   |   |   |   |   |--- type_of_meal_plan_Not Selected <= 0.50
|   |   |   |   |   |   |   |--- weights: [697.09, 9.11] class: 0
|   |   |   |   |   |   |--- type_of_meal_plan_Not Selected >  0.50
|   |   |   |   |   |   |   |--- weights: [15.66, 9.11] class: 0
|   |   |   |   |   |--- lead_time >  102.50
|   |   |   |   |   |   |--- no_of_week_nights <= 2.50
|   |   |   |   |   |   |   |--- weights: [32.06, 19.74] class: 0
|   |   |   |   |   |   |--- no_of_week_nights >  2.50
|   |   |   |   |   |   |   |--- weights: [44.73, 3.04] class: 0
|   |   |   |   |--- market_segment_type_Online >  0.50
|   |   |   |   |   |--- lead_time <= 8.50
|   |   |   |   |   |   |--- lead_time <= 4.50
|   |   |   |   |   |   |   |--- weights: [498.03, 44.03] class: 0
|   |   |   |   |   |   |--- lead_time >  4.50
|   |   |   |   |   |   |   |--- weights: [258.71, 63.76] class: 0
|   |   |   |   |   |--- lead_time >  8.50
|   |   |   |   |   |   |--- required_car_parking_space <= 0.50
|   |   |   |   |   |   |   |--- weights: [2512.51, 1451.32] class: 0
|   |   |   |   |   |   |--- required_car_parking_space >  0.50
|   |   |   |   |   |   |   |--- weights: [134.20, 1.52] class: 0
|   |   |--- no_of_special_requests >  1.50
|   |   |   |--- lead_time <= 90.50
|   |   |   |   |--- no_of_week_nights <= 3.50
```

```
|   |   |   |   |   |--- weights: [1585.04, 0.00] class: 0
|   |   |   |   |--- no_of_week_nights >  3.50
|   |   |   |   |   |--- no_of_special_requests <= 2.50
|   |   |   |   |   |   |--- weights: [180.42, 57.69] class: 0
|   |   |   |   |   |--- no_of_special_requests >  2.50
|   |   |   |   |   |   |--- weights: [52.19, 0.00] class: 0
|   |   |   |--- lead_time >  90.50
|   |   |   |   |--- no_of_special_requests <= 2.50
|   |   |   |   |   |--- arrival_month <= 8.50
|   |   |   |   |   |   |--- weights: [184.90, 56.17] class: 0
|   |   |   |   |   |--- arrival_month >  8.50
|   |   |   |   |   |   |--- weights: [106.61, 106.27] class: 0
|   |   |   |   |--- no_of_special_requests >  2.50
|   |   |   |   |   |--- weights: [67.10, 0.00] class: 0
|--- lead_time >  151.50
|   |--- avg_price_per_room <= 100.04
|   |   |--- no_of_special_requests <= 0.50
|   |   |   |--- no_of_adults <= 1.50
|   |   |   |   |--- market_segment_type_Online <= 0.50
|   |   |   |   |   |--- lead_time <= 163.50
|   |   |   |   |   |   |--- weights: [3.73, 24.29] class: 1
|   |   |   |   |   |--- lead_time >  163.50
|   |   |   |   |   |   |--- weights: [257.96, 62.24] class: 0
|   |   |   |   |--- market_segment_type_Online >  0.50
|   |   |   |   |   |--- avg_price_per_room <= 2.50
|   |   |   |   |   |   |--- weights: [8.95, 3.04] class: 0
|   |   |   |   |   |--- avg_price_per_room >  2.50
|   |   |   |   |   |   |--- weights: [0.75, 97.16] class: 1
|   |   |   |--- no_of_adults >  1.50
|   |   |   |   |--- avg_price_per_room <= 82.47
|   |   |   |   |   |--- market_segment_type_Offline <= 0.50
|   |   |   |   |   |   |--- weights: [2.98, 282.37] class: 1
|   |   |   |   |   |--- market_segment_type_Offline >  0.50
|   |   |   |   |   |   |--- weights: [213.97, 385.60] class: 1
|   |   |   |   |--- avg_price_per_room >  82.47
|   |   |   |   |   |--- no_of_adults <= 2.50
|   |   |   |   |   |   |--- weights: [23.86, 1030.80] class: 1
|   |   |   |   |   |--- no_of_adults >  2.50
|   |   |   |   |   |   |--- weights: [5.22, 0.00] class: 0
|   |   |--- no_of_special_requests >  0.50
|   |   |   |--- no_of_weekend_nights <= 0.50
|   |   |   |   |--- lead_time <= 180.50
|   |   |   |   |   |--- lead_time <= 159.50
|   |   |   |   |   |   |--- weights: [7.46, 7.59] class: 1
|   |   |   |   |   |--- lead_time >  159.50
|   |   |   |   |   |   |--- weights: [37.28, 4.55] class: 0
|   |   |   |   |--- lead_time >  180.50
|   |   |   |   |   |--- no_of_special_requests <= 2.50
|   |   |   |   |   |   |--- weights: [20.13, 212.54] class: 1
|   |   |   |   |   |--- no_of_special_requests >  2.50
```

```
|   |   |   |   |   |   |--- weights: [8.95, 0.00] class: 0
|   |   |   |--- no_of_weekend_nights >  0.50
|   |   |   |   |--- market_segment_type_Offline <= 0.50
|   |   |   |   |   |--- arrival_month <= 11.50
|   |   |   |   |   |   |--- weights: [231.12, 110.82] class: 0
|   |   |   |   |   |--- arrival_month >  11.50
|   |   |   |   |   |   |--- weights: [19.38, 34.92] class: 1
|   |   |   |   |--- market_segment_type_Offline >  0.50
|   |   |   |   |   |--- lead_time <= 348.50
|   |   |   |   |   |   |--- weights: [106.61, 3.04] class: 0
|   |   |   |   |   |--- lead_time >  348.50
|   |   |   |   |   |   |--- weights: [5.96, 4.55] class: 0
|   |--- avg_price_per_room >  100.04
|   |   |--- arrival_month <= 11.50
|   |   |   |--- no_of_special_requests <= 2.50
|   |   |   |   |--- weights: [0.00, 3200.19] class: 1
|   |   |   |--- no_of_special_requests >  2.50
|   |   |   |   |--- weights: [23.11, 0.00] class: 0
|   |   |--- arrival_month >  11.50
|   |   |   |--- no_of_special_requests <= 0.50
|   |   |   |   |--- weights: [35.04, 0.00] class: 0
|   |   |   |--- no_of_special_requests >  0.50
|   |   |   |   |--- arrival_date <= 24.50
|   |   |   |   |   |--- weights: [3.73, 0.00] class: 0
|   |   |   |   |--- arrival_date >  24.50
|   |   |   |   |   |--- weights: [3.73, 22.77] class: 1
```

Observations from decision tree

- We can see that the tree has become simpler and the rules of the trees are readable.
- The model performance of the model has been generalized.
- We observe that the most important features are:
    - Lead Time
    - Market Segment - Online
    - Number of special requests
    - Average price per room

The rules obtained from the decision tree can be interpreted as:

- The rules show that lead time plays a key role in identifying if a booking will be cancelled or not. 151 days has been considered as a threshold value by the model to make the first split.

Bookings made more than 151 days before the date of arrival:

- If the average price per room is greater than 100 euros and the arrival month is December, then the the booking is less likely to be cancelled.
- If the average price per room is less than or equal to 100 euros and the number of special request is 0, then the booking is likely to get canceled.

Bookings made under 151 days before the date of arrival:

- If a customer has at least 1 special request the booking is less likely to be cancelled.
- If the customer didn't make any special requests and the booking was done Online it is more likely to get canceled, if the booking was not done online, it is less likely to be canceled.

If we want more complex then we can go in more depth of the tree

Tuned Model Performance Analysis

The tuned Decision Tree classifier achieved the following results:

Metric	Pre-Pruning Performance
Accuracy	0.83097
Recall	0.78608
Precision	0.72425
F1-score	0.75390

Observations and Insights:

1. **High Accuracy (83.1%)**
    a. The model performs well in classifying both canceled and non-canceled bookings.
    b. However, accuracy alone may not be sufficient, given the class imbalance in the dataset.
2. **Good Recall (78.6%)**
    a. The model correctly identifies most cancellations.
    b. This is beneficial for the hotel since predicting cancellations correctly can help allocate resources efficiently.
3. **Moderate Precision (72.4%)**
    a. The model has some false positives (i.e., predicting a cancellation when the booking is not actually canceled).
    b. A lower precision means that while many cancellations are detected, some predictions might be incorrect.
4. **Balanced F1-Score (75.4%)**
    a. The F1-score is a balance between precision and recall, indicating a reasonable trade-off.
    b. It suggests that the model is neither overly sensitive nor overly specific.

Performance Comparison & Tuning Impact

- Compared to an untuned Decision Tree, this model likely reduces overfitting while maintaining strong predictive performance.
- Pre-pruning methods (such as setting max_depth, min_samples_split) seem to have improved generalization.
- If post-pruning was applied, it may have prevented excessive complexity, leading to stable performance across test data.

Comparison of All Models and Selection of the Best Model

The objective is to compare different models based on their performance metrics and select the best model for predicting hotel booking cancellations.

Performance Comparison of Models

Model	Accuracy	Recall	Precision	F1-score	AUC-ROC
Logistic Regression (Default Threshold = 0.5)	0.80465	0.63089	0.72900	0.67641	0.87
Logistic Regression (Threshold = 0.37)	0.78029	0.79849	0.62994	0.70427	0.87
Logistic Regression (Threshold = 0.42)	0.80992	0.70972	0.71002	0.70987	0.87
Decision Tree (Pre-Pruning)	0.83097	0.78608	0.72425	0.75390	-
Decision Tree (Post-Pruning, if applicable)	TBD	TBD	TBD	TBD	TBD
Other Model (if tested)	TBD	TBD	TBD	TBD	TBD

Analysis of Model Performance

1. Logistic Regression (Default Threshold = 0.5)

- Provides a balanced accuracy (80.5%), recall (63.1%), and precision (72.9%).
- AUC-ROC = 0.87, indicating strong performance in distinguishing between canceled and non-canceled bookings.
- Best suited when probability-based decision-making is needed.

2. Logistic Regression (Threshold = 0.37)

- High recall (79.8%), meaning it identifies most cancellations.
- Lower precision (62.9%), which means more false positives (predicting cancellations that don't happen).
- Use Case: Best when minimizing lost revenue from unfilled rooms due to incorrect cancellations.

3. Logistic Regression (Threshold = 0.42)

- Best trade-off between recall (70.9%) and precision (71.0%).
- F1-score (70.99%) is more balanced compared to the other logistic regression models.
- Use Case: Ideal when the business wants a balance between predicting cancellations accurately and minimizing false positives.

4. Decision Tree (Pre-Pruning)

- Highest accuracy (83.1%), meaning it predicts correctly more often.
- High recall (78.6%), so it detects more cancellations.
- Lower precision (72.4%), meaning it might misclassify some non-cancellations as cancellations.
- Risk of overfitting, even with pre-pruning.
- Use Case: Best when maximizing predictive accuracy is the goal, but needs further validation on generalization.

5. Decision Tree (Post-Pruning)

- Not available yet, but expected to generalize better than the pre-pruned model.
- If available, should be compared to logistic regression for final selection.

Final Recommendation

Best Model: Logistic Regression (Threshold = 0.42)

- Better generalization on new data.
- Good balance between recall (70.9%) and precision (71.0%).
- AUC-ROC = 0.87, showing strong classification ability.
- Easier to interpret, allowing for business insights.
- Can be tuned further with feature engineering.

Model Performance Evaluation and Selection of the Best Model

Based on the models implemented in the Jupyter Notebook (ipynb), we analyze their performance across multiple metrics, including accuracy, recall, precision, F1-score, and AUC-ROC. The goal is to determine which model provides the best balance between predictive accuracy and business needs for predicting hotel booking cancellations.

Performance Comparison of Models

Model	Accuracy	Recall	Precision	F1-score	AUC-ROC
Logistic Regression (Default Threshold = 0.5)	0.80465	0.63089	0.72900	0.67641	0.87
Logistic Regression (Threshold = 0.37)	0.78029	0.79849	0.62994	0.70427	0.87
Logistic Regression (Threshold = 0.42)	0.80992	0.70972	0.71002	0.70987	0.87
Decision Tree (Pre-Pruning)	0.83097	0.78608	0.72425	0.75390	-
Decision Tree (Post-Pruning, if applicable)	TBD	TBD	TBD	TBD	TBD

Model Performance Analysis

1. Logistic Regression (Default Threshold = 0.5)

- Accuracy: 80.46%, indicating a good overall performance.
- Recall: 63.1%, meaning it correctly identifies 63% of cancellations.
- Precision: 72.9%, meaning when it predicts a cancellation, it is correct 72.9% of the time.
- F1-score: 67.6%, reflecting a balance between recall and precision.
- AUC-ROC: 0.87, indicating strong classification ability.

Conclusion:
This model provides a balanced performance but leans slightly towards minimizing false positives (incorrectly predicting a cancellation).

2. Logistic Regression (Threshold = 0.37)

- Accuracy: 78.03% (lower than the default threshold).
- Recall: 79.8% (significantly higher than default).
- Precision: 62.9% (lower than default).
- F1-score: 70.42% (higher than the default model).

Conclusion:

- This threshold increases recall, meaning it captures more cancellations at the cost of slightly lower precision.
- Useful when the goal is to reduce revenue loss from last-minute cancellations

3. Logistic Regression (Threshold = 0.42)

- Accuracy: 80.99% (higher than the default threshold).
- Recall: 70.9% (better than the default model).
- Precision: 71.0% (well-balanced).
- F1-score: 70.99% (most balanced among all logistic regression models).

Conclusion:

- This model provides the best balance between recall and precision.
- It is ideal for business scenarios where predicting cancellations accurately is important without excessive false positives.

4. Decision Tree (Pre-Pruning)

- Accuracy: 83.1% (highest among all models).
- Recall: 78.6% (very high).
- Precision: 72.4% (slightly lower than logistic regression at threshold 0.42).
- F1-score: 75.39% (best among all models).

Conclusion:

- Highly accurate model, but decision trees tend to overfit the training data, which can reduce their reliability on new data.

- If post-pruning is applied and performance remains strong, this model could be a strong candidate.

Best Model: Logistic Regression (Threshold = 0.42)

- Best balance between recall and precision.
- AUC-ROC = 0.87, which is a strong indicator of classification ability.
- Lower risk of overfitting compared to the decision tree.
- Easier to interpret, allowing business stakeholders to extract actionable insights.
- Scalable and computationally efficient, making it suitable for real-time decision-making.

Actionable Insights & Recommendations

Insights

- Overall we can see that the Decision Tree model performs better on the dataset.
- Looking at important variables based on p-values in Logistic regression and feature importance in the Decision Tree model
    - Lead Time, Number of special requests, Average price per room are important in both model
    - From the Logistic Regression model we observe that Lead Time, and Average price per room have a positive relation with bookings getting canclled. And the number of special requests has negative relation with bookings getting cancelled.

Business Recommendations

1. Model Performance & Key Variables Through rigorous analysis, we identify that the Decision Tree model demonstrates superior performance in predicting booking cancellations. By evaluating key variables using p-values in Logistic Regression and feature importance in the Decision Tree model, we observe the following:

Lead Time, Number of Special Requests, and Average Price per Room emerge as critical factors influencing cancellations in both models.

Lead Time and Average Price per Room show a positive correlation with cancellations—longer lead times and higher room prices increase cancellation probability.

Number of Special Requests exhibits a negative correlation—bookings with special requests are less likely to be canceled.

2. Business Strategy Recommendations 2.1. Enhancing Booking Confirmations & Reducing Cancellations Implement automated reminders (emails or notifications) for customers nearing their check-in date to confirm their bookings or make modifications.

Provide pre-arrival reminders about cancellation deadlines and penalties to encourage timely decision-making.

Analyze bookings with high lead times and proactively reach out to customers to assess their likelihood of canceling.

2.2. Optimizing Cancellation Policies for Profitability Stricter refund policies should be applied to high-value bookings, particularly those with special requests, to mitigate financial losses.

Online bookings, which show a higher cancellation rate, should be subject to lower refund percentages to discourage last-minute cancellations.

Transparent cancellation policies should be clearly displayed on hotel websites and booking platforms to ensure customer awareness.

2.3. Managing Length of Stay to Reduce Cancellations Data suggests that bookings exceeding five days have a higher probability of cancellation.

Implement a two-phase booking process: Allow initial reservations for up to five days, with an option to extend upon reconfirmation.

This policy can be selectively applied—corporate and aviation market segments may be given more flexibility, while leisure travelers may require re-confirmation for longer stays.

2.4. Seasonal Strategy for Resource Allocation December and January have a lower cancellation ratio—resources should be optimized to handle peak demand.

October and September see the highest number of bookings and cancellations—hotels should investigate the root causes (e.g., promotions, event-driven demand).

2.5. Strengthening Customer Engagement & Loyalty Post-booking interactions can enhance customer experience and reduce cancellations by reinforcing commitment.

Provide personalized recommendations (local events, attractions, dining options).

Offer exclusive perks for early confirmations.

Loyalty Program for Repeat Guests

Given that repeat guests have a significantly lower cancellation rate, hotels should invest in retaining them.

Introduce personalized discounts, priority check-ins, and exclusive services to encourage repeat bookings.

A well-structured loyalty program fosters long-term relationships and enhances customer lifetime value.

Conclusion By leveraging predictive analytics, hotels can minimize cancellations, optimize resource allocation, and enhance customer retention. A hybrid approach, combining personalized engagement strategies, dynamic pricing models, and revised cancellation policies, can significantly improve operational efficiency and guest satisfaction.